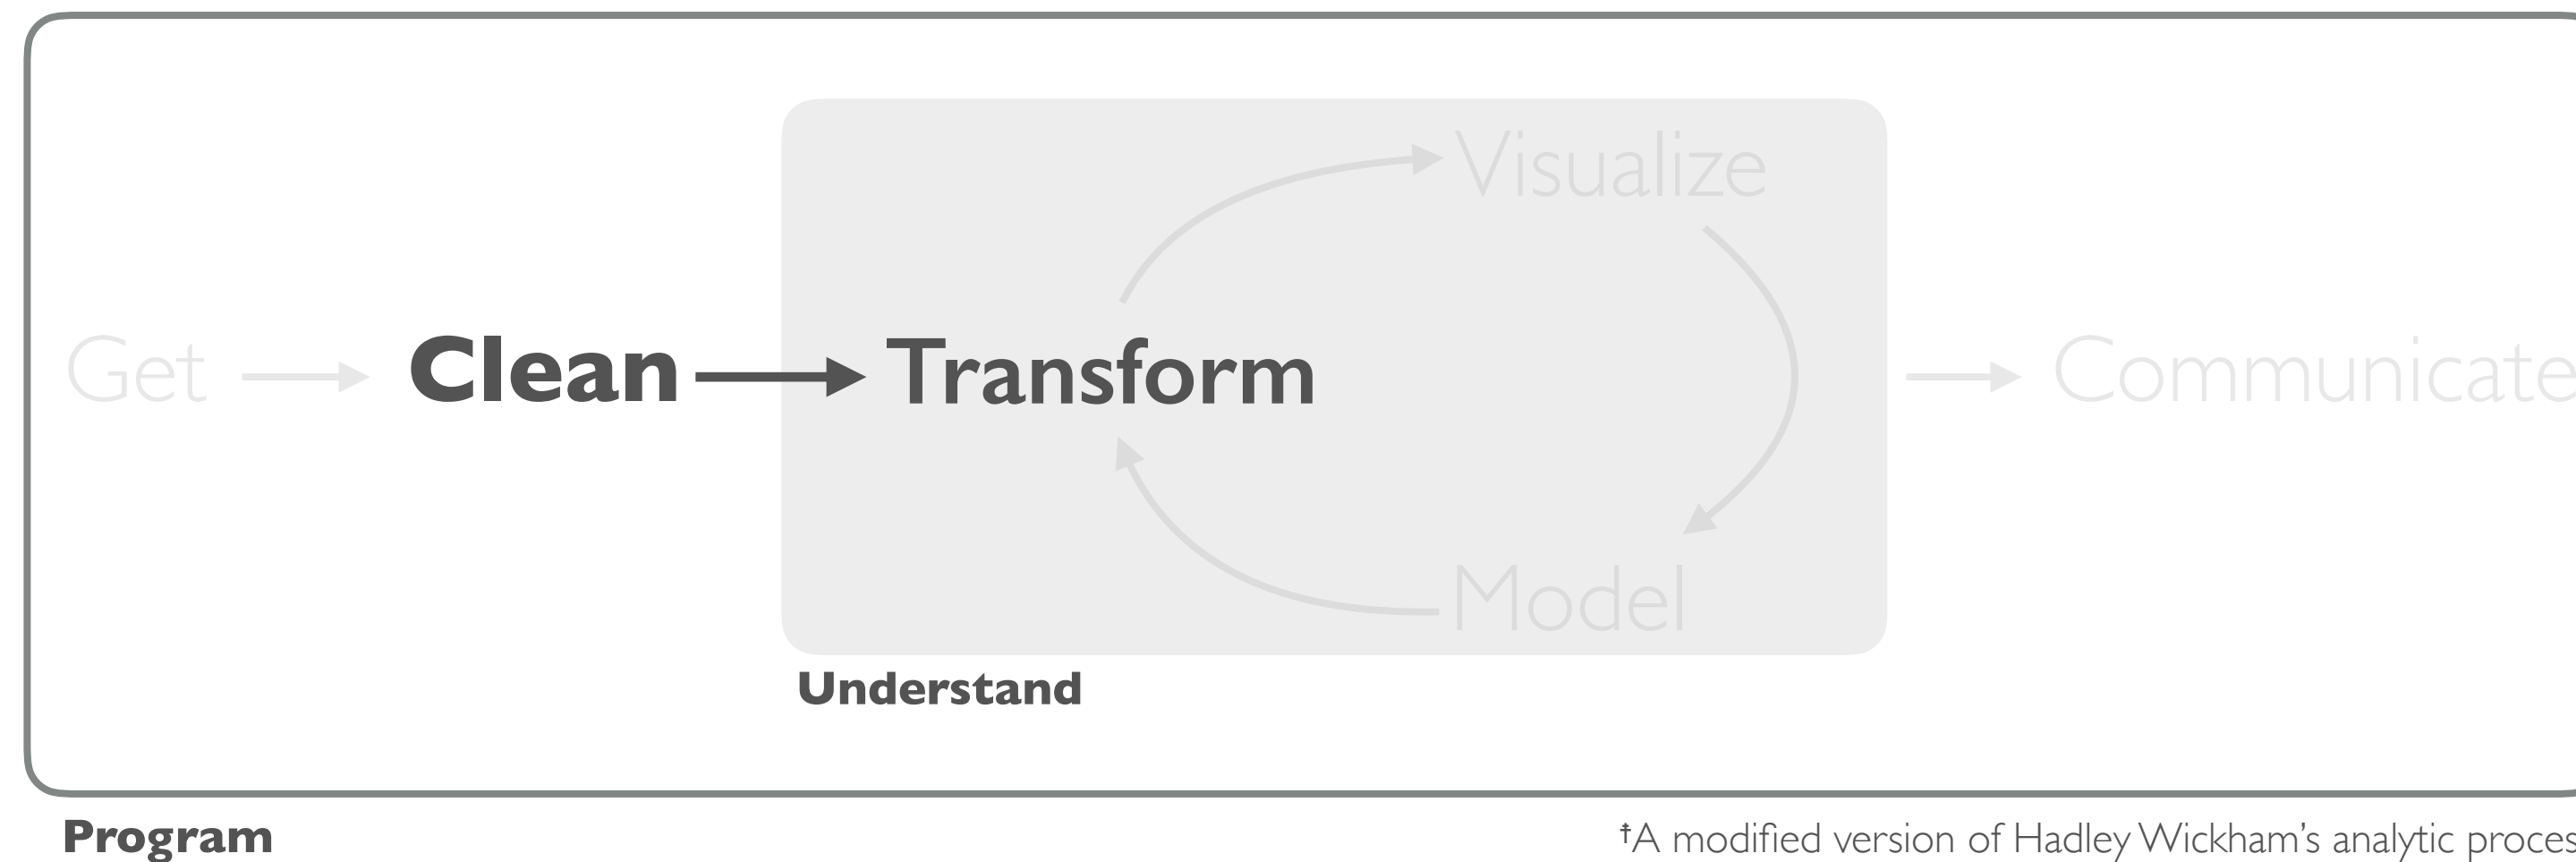



# REGULAR EXPRESSIONS



“Analysts are often trained to handle tabular or rectangular data that are mostly numeric, but much of the data proliferating today are unstructured and text-heavy.”

— Julia Silge and David Robinson

# PREREQUISITES



# PACKAGE PREREQUISITE

```
library(tidyverse)
```

```
if (packageVersion("devtools") < 1.6) {  
  install.packages("devtools")  
}
```

```
devtools::install_github("bradleyboehmke/harrypotter")  
library(harrypotter)
```

# REGULAR EXPRESSIONS

- Regular expressions (regex) are strings to identify patterns in text
- Two areas of focus:
  - i. regex functions
  - ii. regex syntax
- The **stringr** package provides us a convenient approach to regex text mining
- We'll explore dealing with regex in both character strings and data frames



# REGEX FUNCTIONS

Dealing with character strings



# DATA PREREQUISITE

philosophers\_stone

[1] "THE BOY WHO LIVED     Mr. and Mrs. Dursley, of number four, Privet Drive, were proud to say that they were perfectly normal, thank you very much. They were the last people you'd expect to be involved in anything strange or mysterious, because they just didn't hold with such nonsense.     Mr. Dursley was the director of a firm called Grunnings, which made drills. He was a big, beefy man with hardly any neck, although he did have a very large mustache. Mrs. Dursley was thin and blonde and had nearly twice the usual amount of neck, which came in very useful as she spent so much of her time craning over garden fences, spying on the neighbors. The Dursleys had a small son called Dudley and in their opinion there was no finer boy anywhere.     The Dursleys had everything they wanted, but they also had a secret, and their greatest fear was that somebody would discover it. They didn't think they could bear it if anyone found out about the Potters. Mrs. Potter was Mrs. Dursley's sister, but they h... <truncated>

[2] "THE VANISHING GLASS     Nearly ten years had passed since the Dursleys had woken up to find their nephew on the front step, but Privet Drive had hardly changed at all. The sun rose on the same tidy front gardens and lit up the brass number four on the Dursleys' front door; it crept into their living room, which was almost exactly the same as it had

# str\_\*

```
str_*(string, pattern)
```

**string**: character vector

**pattern**: regex pattern to look for



# str\_\*

```
str_detect(philosophers_stone, "Harry")  
[1] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE  
[10] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
```

**str\_detect:** does the expression exist?

# str\_\*

```
str_detect(philosophers_stone, "Harry")
```

```
str_count(philosophers_stone, "Harry")
```

```
[1] 20 79 68 48 147 117 63 49 70 82 68  
[12] 99 62 45 93 116 99
```

**str\_count:** how many instances  
are there?

# str\_\*

```
str_detect(philosophers_stone, "Harry")
```

```
str_count(philosophers_stone, "Harry")
```

```
str_extract(philosophers_stone, "Harry")
```

```
[1] "Harry" "Harry" "Harry" "Harry" "Harry"
```

```
[6] "Harry" "Harry" "Harry" "Harry" "Harry"
```

```
[11] "Harry" "Harry" "Harry" "Harry" "Harry"
```

```
[16] "Harry" "Harry"
```

`str_extract`: extract the **first** instance

# str\_\*

```
str_detect(philosophers_stone, "Harry")
```

```
str_count(philosophers_stone, "Harry")
```

```
str_extract(philosophers_stone, "Harry")
```

```
str_extract_all(philosophers_stone, "Harry")
```

```
[[1]]
```

```
[1] "Harry" "Harry" "Harry" "Harry" "Harry"
```

```
[6] "Harry" "Harry" "Harry" "Harry" "Harry"
```

```
[11] "Harry" "Harry" "Harry" "Harry" "Harry"
```

```
[16] "Harry" "Harry" "Harry" "Harry" "Harry"
```

```
[[2]]
```

```
[1] "Harry" "Harry" "Harry" "Harry" "Harry"
```

```
[6] "Harry" "Harry" "Harry" "Harry" "Harry"
```

```
[11] "Harry" "Harry" "Harry" "Harry" "Harry"
```

```
[16] "Harry" "Harry" "Harry" "Harry" "Harry"
```

`str_extract_all`: extract **all** instances

# str\_\*

```
str_detect(philosophers_stone, "Harry")
str_count(philosophers_stone, "Harry")
str_extract(philosophers_stone, "Harry")
str_extract_all(philosophers_stone, "Harry")
```

```
str_locate_all(philosophers_stone, "Harry")
```

```
[[1]]
```

	start	end
[1,]	5243	5247
[2,]	5798	5802
[3,]	5868	5872
[4,]	10231	10235
[5,]	18057	18061
[6,]	18190	18194
[7,]	18521	18525

`str_locate_all`: locate the position of **all** instances

# YOUR TURN!

*Take 5 minutes to explore the various **str\_\*** functions*

# REGEX SYNTAX

Dealing with character strings

`([a-z][^a-z0-9`

# MULTIPLE WORDS / CASE SENSITIVE

```
str_count(philosophers_stone, "Harry Potter")
```

```
[1] 5 2 0 2 3 5 1 1 0 0 0 0 0 1 5 0 3
```

Phrases



# MULTIPLE WORDS / CASE SENSITIVE

```
str_count(philosophers_stone, "Harry Potter")
```

```
str_count(philosophers_stone, "Harry|Potter")
```

```
[1] 28 57 56 30 109 76 50 44 55 67 52 73 52
```

```
[14] 31 70 80 61
```

“Harry” or “Potter”

# MULTIPLE WORDS / CASE SENSITIVE

```
str_count(philosophers_stone, "Harry Potter")
```

```
str_count(philosophers_stone, "Harry I Potter")
```

```
str_count(philosophers_stone, "ye(slah)")
```

```
[1] 17  5  3  7 25  9 13  5  8  4  5 10  1  4  8 12  8
```

“yes” or “yeah”

# MULTIPLE WORDS / CASE SENSITIVE

```
str_count(philosophers_stone, "Harry Potter")
```

```
str_count(philosophers_stone, "Harry I Potter")
```

```
str_count(philosophers_stone, "ye(slah)")
```

```
str_count(philosophers_stone, "boy")
```

```
[1]  9  7  2  5 15 34  5  1  7  1  0  3  0  0  3  2  6
```

← default

```
str_count(philosophers_stone, regex("boy", ignore_case = TRUE))
```

```
[1] 10  7  2  5 15 34  5  1  7  1  0  3  0  0  3  2  6
```

← ignore case

# YOUR TURN!

*How many times are “Mr” and “Mrs” used in  
philosophers\_stone?*

# SOLUTION

```
str_count(philosophers_stone, "Mr\\.\\.s|Mrs\\.\\.s") %>% sum  
[1] 126
```

# ANCHORS

```
str_count(deathly_hallows, "^Harry")
[1] 0 1 0 1 0 0 0 0 0 1 0 0 0 1 0 0 1 0 0 1 1 1 1 0 0 0 0
[28] 1 0 0 0 0 1 0 0 0 0
```

^: Identify patterns at the **beginning** of an element

# ANCHORS

```
str_count(deathly_hallows, "^Harry")
str_count(philosophers_stone,
          regex("end$", ignore_case = TRUE))
[1] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1
```

**\$**: Identify patterns at the **end** of an element

# YOUR TURN!

*Extract all elements in **deathly\_hallows** that start with  
“Harry”*



# SOLUTION

```
deathly_hallows[str_detect(deathly_hallows, "^Harry")]
```

[1] "Harry was bleeding. Clutching his right hand in his left and swearing under his breath, he shouldered open his bedroom door. There was a crunch of breaking china. He had trodden on a cup of cold tea that had been sitting on the floor outside his bedroom door. \"What the \"C?\"He looked around, the landing of number four, Privet Drive, was deserted. Possibly the cup of tea was Dudley's idea of a clever booby trap. Keeping his bleeding hand elevated, Harry scraped the fragments of cup together with the other hand and threw them into the already crammed bin just visible inside his bedroom door. Then he tramped across to the bathroom to run his finger under the tap. It was stupid, pointless, irritating beyond belief that he still had four days left of being unable to perform magic...but he had to admit to himself that this jagged cut in his finger would have defeated him. He had never learned how to repair wounds, and now he came to think of it \"C particularly in light of his immediate ... <truncated>

[2] "Harry ran back upstairs to his bedroom, arriving at the window just in time to see the Dursleys' car swinging out of the drive and off up the road. Dedalus's top hat was visible between Aunt Petunia and Dudley in the backseat. The car turned right at the end of Privet Drive, its windows burned scarlet for a moment in the now setting sun, and then it was gone. Harry picked up Hedwig's cage, his Firebolt, and his rucksack, gave his

# SPECIAL PATTERNS

```
str_extract(philosophers_stone, "Harry.")  
[1] "Harry\" " "Harry \" " "Harry \" " "Harry!" "  
[5] "Harry \" " "Harry' \" " "Harry' \" " "Harry \" "  
[9] "Harry \" " "Harry \" " "Harry \" " "Harry \" "  
[13] "Harry \" " "Harry, \" " "Harry' \" " "Harry \" "  
[17] "Harry.\" "
```

. : wild card - **any character**

# SPECIAL PATTERNS

```
str_extract(philosophers_stone, "Harry.")
```

```
str_extract(philosophers_stone, "\\d")
```

```
[1] NA NA "4" "1" "0" "1" NA "3" "3" NA "1" NA  
[13] "1" "1" NA "1" NA
```

**\\d**: digits

# SPECIAL PATTERNS

```
str_detect(philosophers_stone, "Harry")  
str_extract(philosophers_stone, "\\d")  
str_extract(philosophers_stone, "[1|4]")  
[1] NA NA "4" "1" "1" "1" NA "1" NA NA "1" NA  
[13] "1" "1" NA "1" NA
```

`[d|d]`: specified digits

# SPECIAL PATTERNS

```
str_detect(philosophers_stone, "Harry")
```

```
str_count(philosophers_stone, "Harry")
```

```
str_extract(philosophers_stone, "[1|4]")
```

```
str_extract(philosophers_stone, "[yz].")
```

```
[1] "ey," "ly" "azi" "ey" "ry" "ry'" "ry" " yo"
```

```
[9] "ry" "oy" "ey" "oze" "ry" "ey'" "dy" " ye"
```

```
[17] "ry."
```

**[a-zA-Z]**: specified letters

# YOUR TURN!

*How many times is the word “Harry” get followed by a word that starts with a vowel in **philosophers\_stone**?*

# SOLUTION

```
str_count(philosophers_stone, "Harry.[aeiou]") %>% sum  
str_count(philosophers_stone, "Harry\\s[aeiou]") %>% sum  
[1] 156
```



This is the safer option since `\\s` specifically identifies white space

# REPETITION

```
str_extract(philosophers_stone, "[aeiou]{4}")  
str_extract(philosophers_stone, "[aeiou]{3,}")  
str_extract(philosophers_stone, "[aeiou]{3,4}")  
[1] "iou" "uie" "iou" "uee" "iou" "iou" "uie" "uea"  
[9] "uie" "iou" "iou" "iou" "eei" "iou" "iou" "iou"  
[17] "uie"
```

**{n}**: find n repetitions

**{n,}**: find n or more repetitions

**{n,m}**: find n to m repetitions



# YOUR TURN!

1. *Without computer support, what is this finding:*

```
str_count(philosophers_stone, regex("(no[[:punct:]])[ ]){3}", ignore_case = TRUE))
```

2. *Extract the 25 characters that precede and follow the use of “Harry” in **philosophers\_stone***

# SOLUTION

```
str_count(philosophers_stone, regex("(no[[:punct:]])[ ]){3}", ignore_case = TRUE))
```

Looking for 3 repeated uses of “no” with any form of punctuation and a space following each use of “no”. (i.e. “No, no, no!”)

```
str_extract_all(philosophers_stone, “.{25}Harry.{25}”)
```

```
[1] " I heard yes, their son, Harry\" Mr. Dursley stopped de"  
[2] "ter who had a son called Harry. Come to think of it, he"  
[3] "re his nephew was called Harry. He'd never even seen th"  
[4] "in? Howard, isn't it?\" \"Harry. Nasty, common name, if "  
[5] "o kill the Potter's son, Harry. But -- he couldn't. He "  
[6] "at when he couldn't kill Harry Potter, Voldemort's powe"  
[7] "n the name of heaven did Harry survive?\" \"We can only "  
[8] "s?\" \"I've come to bring Harry to his aunt and uncle. T"  
[9] "t, screaming for sweets. Harry Potter come and live her"  
[10] "ed if today was known as Harry Potter day in the future"  
[11] "l, he books written about Harry -- every child in our wo"
```

# REGEX

Doing similar stuff with a data frame

*fix*

# DATA PREREQUISITE

```
airbnb <- read_rds("data/airbnb.rds")
```

airbnb

# A tibble: 3,585 × 95

	id	listing_url	scrape_id	last_scraped
	<int>	<chr>	<dbl>	<date>
1	12147973	https://www.airbnb.com/rooms/12147973	2.016091e+13	2016-09-07
2	3075044	https://www.airbnb.com/rooms/3075044	2.016091e+13	2016-09-07
3	6976	https://www.airbnb.com/rooms/6976	2.016091e+13	2016-09-07
4	1436513	https://www.airbnb.com/rooms/1436513	2.016091e+13	2016-09-07
5	7651065	https://www.airbnb.com/rooms/7651065	2.016091e+13	2016-09-07
6	12386020	https://www.airbnb.com/rooms/12386020	2.016091e+13	2016-09-07
7	5706985	https://www.airbnb.com/rooms/5706985	2.016091e+13	2016-09-07
8	2843445	https://www.airbnb.com/rooms/2843445	2.016091e+13	2016-09-07
9	753446	https://www.airbnb.com/rooms/753446	2.016091e+13	2016-09-07
10	840408	https://www.airbnb.com/rooms/840408	2.016091e+13	2016-09-07

# DATA FRAME REGEX MADE EASY



+



=



# BASICS

```
airbnb %>%  
  select(name) %>%  
  mutate(character_count = str_count(name))  
# A tibble: 3,585 × 2
```

	name	character_count
	<chr>	<int>
1	Sunny Bungalow in the City	26
2	Charming room in pet friendly apt	33
3	Mexican Folk Art Haven in Boston	32
4	Spacious Sunny Bedroom Suite in Historic Home	45
5	Come Home to Boston	19
6	Private Bedroom + Great Coffee	30
7	New Lrg Studio apt 15 min to Boston	35
8	"Tranquility" on "Top of the Hill"	34
9	6 miles away from downtown Boston!	34
10	Perfect & Practical Boston Rental	33

```
# ... with 3,575 more rows
```

- We can use **str\_count** to count the number of characters in a character field



# BASICS

```
airbnb %>%
  select(name) %>%
  mutate(first_five = str_sub(name, start = 1, end = 5),
         last_five = str_sub(name, start = -5))
# A tibble: 3,585 × 3
```

	name	first_five	last_five
	<chr>	<chr>	<chr>
1	Sunny Bungalow in the City	Sunny	City
2	Charming room in pet friendly apt	Charm	y apt
3	Mexican Folk Art Haven in Boston	Mexic	oston
4	Spacious Sunny Bedroom Suite in Historic Home	Spaci	Home
5	Come Home to Boston	Come	oston
6	Private Bedroom + Great Coffee	Priva	offee
7	New Lrg Studio apt 15 min to Boston	New L	oston
8	"Tranquility" on "Top of the Hill"	"Tran	Hill"
9	6 miles away from downtown Boston!	6 mil	ston!
10	Perfect & Practical Boston Rental	Perfe	ental

```
# ... with 3,575 more rows
```

- We can use `str_sub` with `start` and `end` arguments to take out a substring

# BASICS

```
airbnb %>%
  select(host_name) %>%
  mutate(lower_case = str_to_lower(host_name),
         upper_case = str_to_upper(host_name))
# A tibble: 3,585 × 3
  host_name lower_case upper_case
  <chr>      <chr>      <chr>
1 Virginia  virginia    VIRGINIA
2 Andrea    andrea      ANDREA
3 Phil      phil        PHIL
4 Meghna    meghna      MEGHNA
5 Linda     linda       LINDA
6 Deborah   deborah     DEBORAH
7 Juliet    juliet      JULIET
8 Marilyn   marilyn     MARILYN
9 Sami      sami        SAMI
10 Damon    damon       DAMON
# ... with 3,575 more rows
```

- We can use `str_to_lower` and `str_to_upper` to normalize text case



# YOUR TURN!

1. *What is the average number of characters used in the **name** column? What about the **description** column?*
2. *What is the most common name in the **host\_name** column?*

# SOLUTION

```
# problem 1
airbnb %>%
  select(name, description) %>%
  mutate(
    name_char = str_count(name),
    desc_char = str_count(description)
  ) %>%
  summarise(
    name_char = mean(name_char, na.rm = TRUE),
    desc_char = mean(desc_char, na.rm = TRUE)
  )
# A tibble: 1 × 2
  name_char desc_char
    <dbl>     <dbl>
1  32.34728  768.5431
```

# SOLUTION

```
# problem 2
airbnb %>%
  select(host_name) %>%
  mutate(host_name = str_to_lower(host_name)) %>%
  count(host_name, sort = TRUE)
# A tibble: 1,334 × 2
  host_name      n
  <chr> <int>
1      kara    138
2  seamless    79
3      mike    71
4  flatbook    58
5    alicia    50
6    marie    42
7    jason    35
8    sarah    26
```

# FILTERING

```
airbnb %>%
  select(name) %>%
  mutate(charming = str_detect(name, regex("charming", ignore_case = TRUE)))
# A tibble: 3,585 × 2
```

	name	charming
	<chr>	<lgl>
1	Sunny Bungalow in the City	FALSE
2	Charming room in pet friendly apt	TRUE
3	Mexican Folk Art Haven in Boston	FALSE
4	Spacious Sunny Bedroom Suite in Historic Home	FALSE
5	Come Home to Boston	FALSE
6	Private Bedroom + Great Coffee	FALSE
7	New Lrg Studio apt 15 min to Boston	FALSE
8	"Tranquility" on "Top of the Hill"	FALSE
9	6 miles away from downtown Boston!	FALSE
10	Perfect & Practical Boston Rental	FALSE

```
# ... with 3,575 more rows
```

- We can use `str_detect` to see if the word “charming” exists in the name
- Since `str_detect` supplies a logical response we can use this for filtering...

# FILTERING

```
airbnb %>%
  select(name) %>%
  filter(str_detect(name, regex("charming", ignore_case = TRUE)))
# A tibble: 92 x 1
#   name
#   <chr>
1 Charming room in pet friendly apt
2 Cozy room in a charming villa.
3 Charming Gambrel on a sweet street
4 Charming 3 bedroom-15 min to Boston
5 Charming new house-15 min to Boston
6 Queen room in a charming villa
7 Charming sunlit house in Boston
8 Charming Victorian near T
9 Charming 2BD Across from Arboretum
10 Charming Boston Apartment
# ... with 82 more rows
```

- We can use `str_detect` to see if the word “charming” exists in the name
- Since `str_detect` supplies a logical response we can use this for filtering...

# FILTERING

- We can use different approaches to get to the same results
- What do you expect these to return????

```
airbnb %>%  
  select(name) %>%  
  filter(str_detect(name, "(C|c)harming|(C|c)ute"))
```

```
airbnb %>%  
  select(name) %>%  
  mutate(name = str_to_lower(name)) %>%  
  filter(str_detect(name, "charming|cute"))
```

# FILTERING

- We can use different approaches to get to the same results

```
airbnb %>%
  select(name) %>%
  filter(str_detect(name, "(C|c)harming|(C|c)ute"))
# A tibble: 105 x 1
  name
  <chr>
1 Charming room in pet friendly apt
2 Cozy room in a charming villa.
3 Charming Gambrel on a sweet street
4 Charming 3 bedroom-15 min to Boston
5 Charming new house-15 min to Boston
6 Queen room in a charming villa
7 Charming sunlit house in Boston
8 Charming Victorian near T
9 Charming 2BD Across from Arboretum
```

```
airbnb %>%
  select(name) %>%
  mutate(name = str_to_lower(name)) %>%
  filter(str_detect(name, "charming|cute"))
# A tibble: 105 x 1
  name
  <chr>
1 Charming room in pet friendly apt
2 Cozy room in a charming villa.
3 Charming Gambrel on a sweet street
4 Charming 3 bedroom-15 min to Boston
5 Charming new house-15 min to Boston
6 Queen room in a charming villa
7 Charming sunlit house in Boston
8 Charming Victorian near T
```

# YOUR TURN!

1. Using the **house\_rules** column, how many observations (aka hosts) advocate for “no shoes”?
2. How would you filter out these observations?



# SOLUTION

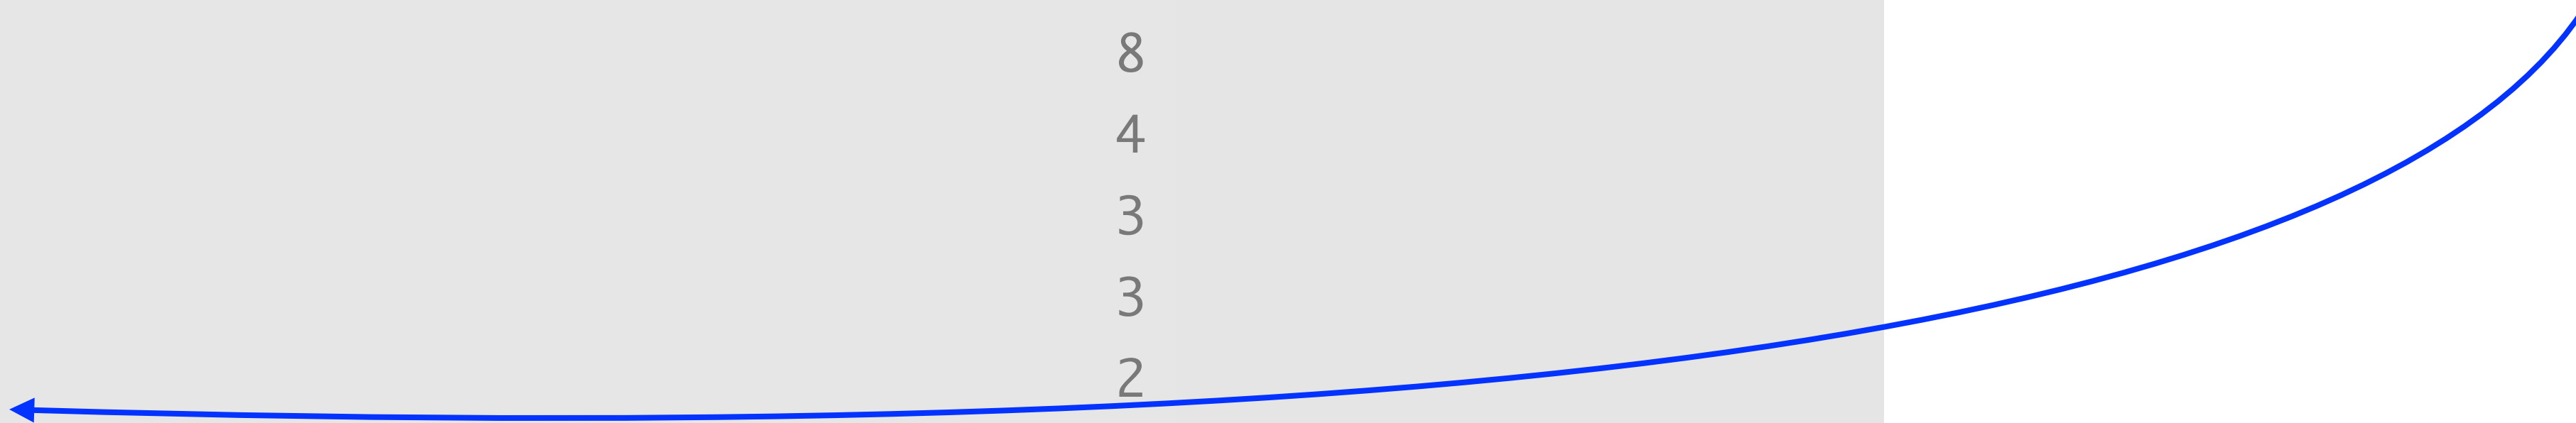
```
# problem 1
airbnb %>%
  tally(str_detect(house_rules, regex("no shoes", ignore_case = TRUE)))
# A tibble: 1 × 1
  no_shoes
    <int>
1         67

# problem 2
airbnb %>%
  filter(!str_detect(house_rules, regex("no shoes", ignore_case = TRUE)))
# A tibble: 2,326 × 95
   id listing_url scrape_id last_scraped
  <int>         <chr>      <dbl>      <date>
1 12147973 https://www.airbnb.com/rooms/12147973 2.016091e+13 2016-09-07
2  3075044 https://www.airbnb.com/rooms/3075044 2.016091e+13 2016-09-07
3    6976 https://www.airbnb.com/rooms/6976 2.016091e+13 2016-09-07
4 1126512 https://www.airbnb.com/rooms/1126512 2.016091e+13 2016-09-07
```

# CLEANING

```
airbnb %>%  
  select(name) %>%  
  count(str_extract(name, "^[^A-Za-z0-9]+"), sort = TRUE)  
# A tibble: 20 x 2  
  `str_extract(name, "^[^A-Za-z0-9]+")`      n  
  <chr>                                <int>  
1 NA                                  3468  
2 [                                   79  
3 $                                    8  
4 "\\\" 4  
5 (                                    3  
6 #                                    3  
7 *                                    2  
8 **                                   2  
9 ^                                    2  
10 【                                  1  
11 "* "                                1  
12 "*** "                              1
```

- Sometimes we need to do some cleaning. For example, if we wanted to look for the most common first words used in names, we may want to clean up non-alphanumeric characters.



# CLEANING

```
airbnb %>%
  select(name) %>%
  mutate(
    name = str_replace_all(name, "[^A-Za-z0-9]+", " "),
    name = str_replace_all(name, "[[:punct:]]+", " "),
    name = str_trim(name),
    name = str_to_lower(name)
  ) %>%
  count(str_extract(name, "^[A-Za-z0-9]+"), sort = TRUE)
```

# A tibble: 626 x 2

	`str_extract(name, "^[A-Za-z0-9]+")`	n
	<chr>	<int>
1	cozy	183
2	private	169
3	beautiful	120
4	spacious	114
5	lux	104

- Remove all non-alphanumerics
- Remove punctuations
- Remove extra white spaces
- Standardize to lowercase
- Extract and count first alphanumeric words

CHALLENGE



# CHALLENGE

*In the Kaggle competition for predicting Titanic survivors, the most important predictor variable ended up being the passenger's title (i.e. Mr., Mrs., Miss., Master).*

*Using the `titanic::titanic_train` data, extract the passengers title and create a new feature named "Title".*

# SOLUTION

```
titanic <- titanic::titanic_train %>% as_tibble()
```

```
titanic %>%  
  mutate(Title = str_replace_all(Name, "(.*, )|(\\"..*)", ""))
```


# A tibble: 891 x 13

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked	Title
	<int>	<int>	<int>	<chr>	<chr>	<dbl>	<int>	<int>	<chr>	<dbl>	<chr>	<chr>	<chr>
1	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.25	""	S	Mr
2	2	1	1	Cumings, Mrs. John Bradley (Floren...	female	38.0	1	0	PC 17599	71.3	C85	C	Mrs
3	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 31...	7.92	""	S	Miss
4	4	1	1	Futrelle, Mrs. Jacques Heath (Lily...	female	35.0	1	0	113803	53.1	C123	S	Mrs
5	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.05	""	S	Mr
6	6	0	3	Moran, Mr. James	male	NA	0	0	330877	8.46	""	Q	Mr
7	7	0	1	McCarthy, Mr. Timothy J	male	54.0	0	0	17463	51.9	E46	S	Mr
8	8	0	3	Palsson, Master. Gosta Leonard	male	2.00	3	1	349909	21.1	""	S	Mast...
9	9	1	3	Johnson, Mrs. Oscar W (Elisabeth V...	female	27.0	0	2	347742	11.1	""	S	Mrs
10	10	1	2	Nasser, Mrs. Nicholas (Adele Achem)	female	14.0	1	0	237736	30.1	""	C	Mrs



# ... with 881 more rows



SO LITTLE TIME!



# LEARN MORE





WHAT TO REMEMBER



# FUNCTIONS TO REMEMBER

Operator/Function	Description
<code>str_*</code>	stringr functions for regular expressions
<code>regex(pattern, ignore_case = TRUE)</code>	ignore case
<code>“x   y” “(x y)”</code>	using or for finding multiple forms of regular expressions
<code>^ \$</code>	anchors - finding regex at beginning or end of element
<code>., \\d, \\s, [0-9], [a-zA-Z]</code>	finding regex patterns
<code>+, {n}, {n,}, {n,m}</code>	finding repetitions of regex patterns