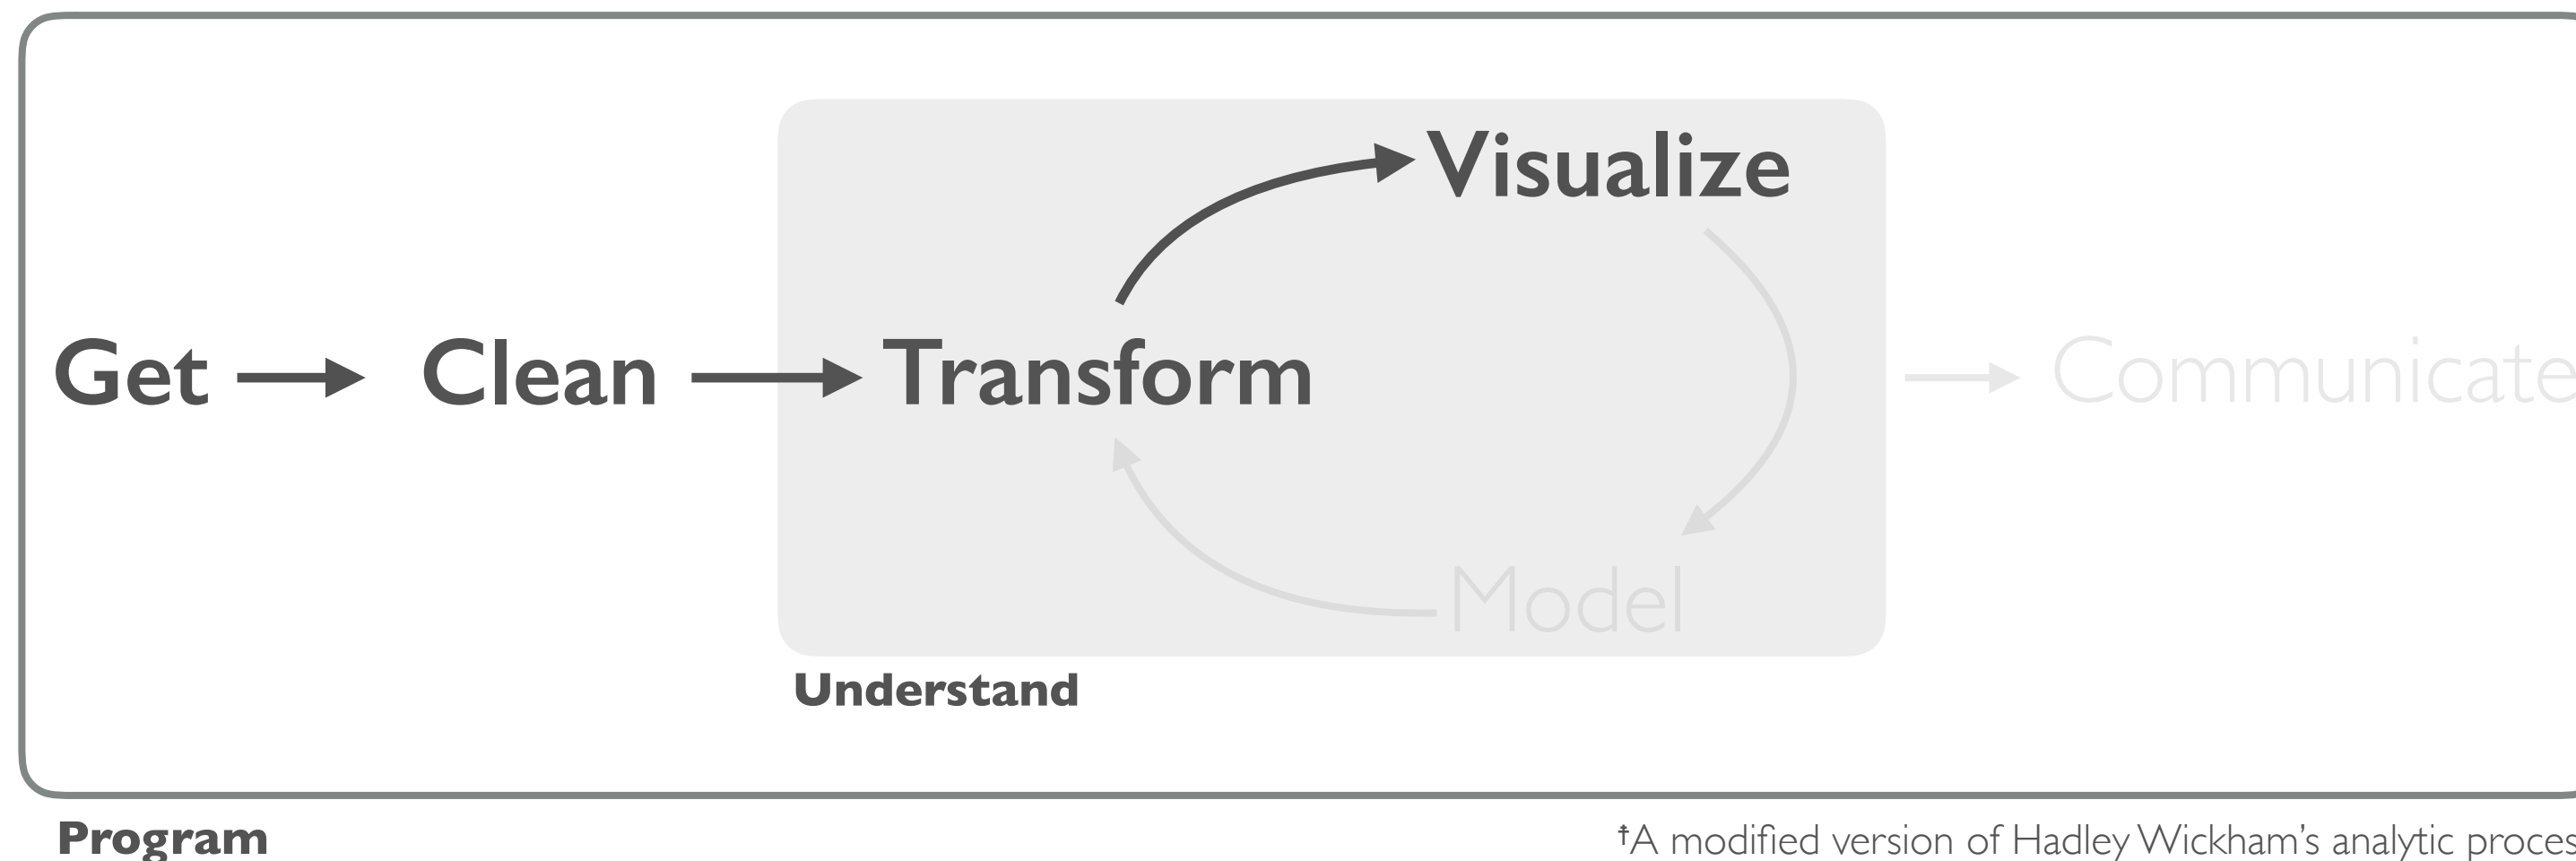


ORGANIZE, CLEAN, & ANALYZE



†A modified version of Hadley Wickham's analytic process

“Using tidy data principles is a powerful way to make handling data easier and more effective, and this is no less true when it comes to dealing with text.”

— Julia Silge and David Robinson

WHAT IS TIDY DATA?

- One variable per column
- One observation per row

##		userid	age	dob_day	dob_year	dob_month	gender	tenure	friend_count
##	1	2094382	14	19	1999	11	male	266	0
##	2	1192601	14	2	1999	11	female	6	0
##	3	2083884	14	16	1999	11	male	13	0
##	4	1203168	14	25	1999	12	female	93	0
##	5	1733186	14	4	1999	12	male	82	0
##	6	1524765	14	1	1999	12	male	15	0
##	7	1136133	13	14	2000	1	male	12	0
##	8	1680361	13	4	2000	1	female	0	0
##	9	1365174	13	1	2000	1	male	81	0
##	10	1712567	13	2	2000	2	male	171	0
##	11	1612453	13	22	2000	2	male	98	0
##	12	2104073	13	1	2000	2	male	55	0

TIDY TEXT

From this...

philosophers_stone


[1] "THE BOY WHO LIVED Mr. and Mrs. Dursley, of number four, Privet Drive, were proud to say that they were perfectly normal, thank you very much. They were the last people you'd expect to be involved in anything strange or mysterious, because they just didn't hold with such nonsense. Mr. Dursley was the director of a firm called Grunnings, which made drills. He was a big, beefy man with hardly any neck, although he did have a very large mustache. Mrs. Dursley was thin and blonde and had nearly twice the usual amount of neck, which came in very useful as she spent so much of her time craning over garden fences, spying on the neighbors. The Dursleys had a small son called Dudley and in their opinion there was no finer boy anywhere. The Dursleys had everything they wanted, but they also had a secret, and their greatest fear was that somebody would discover it. They didn't think they could bear it if anyone found out about the Potters. Mrs. Potter was Mrs. Dursley's sister, but they h...
<truncated>

[2] "THE VANISHING GLASS Nearly ten years had passed since the Dursleys had woken up to find their nephew on the front

To this...

book	chapter	word
<chr>	<int>	<chr>
1 Philosopher's Stone	1	boy
2 Philosopher's Stone	1	lived
3 Philosopher's Stone	1	dursley
4 Philosopher's Stone	1	privet
5 Philosopher's Stone	1	drive
6 Philosopher's Stone	1	proud
7 Philosopher's Stone	1	perfectly
8 Philosopher's Stone	1	normal
9 Philosopher's Stone	1	people
10 Philosopher's Stone	1	expect
# ... with 409,328 more rows		


TIDY TEXT



STEPS

1. Get text into a data frame
2. Unnest to specified token size
3. Remove stop words
4. Perform frequency analysis
5. Compare proportions across documents

PREREQUISITES



PACKAGE PREREQUISITE

```
library(tidyverse)
```

```
library(tidytext)
```

```
if (packageVersion("devtools") < 1.6) {  
  install.packages("devtools")  
}
```

```
devtools::install_github("bradleyboehmke/harrypotter")
```

```
library(harrypotter)
```


CREATING STRUCTURE

Getting text into a data frame



TEXT TO TIBBLE

- Often text is held in a vector of character strings

```
text <- c("Because I could not stop for Death -",  
         "He kindly stopped for me -",  
         "The Carriage held but just Ourselves -",  
         "and Immortality")
```

```
text  
[1] "Because I could not stop for Death -" "He kindly stopped for me -"  
[3] "The Carriage held but just Ourselves -" "and Immortality"
```

TEXT TO TIBBLE

- Often text is held in a vector of character strings
- We can convert these into a data frame (tibble) quite easily

```
tibble(  
  index = seq_along(text),  
  token = text  
)  
# A tibble: 4 x 2  
  index token  
  <int> <chr>  
1     1 1 Because I could not stop for Death -  
2     2 2 He kindly stopped for me -  
3     3 3 The Carriage held but just Ourselves -  
4     4 4 and Immortality
```

Creates a tibble with two variables

- chapter
- all text for that chapter

DATA PREREQUISITE

philosophers_stone

[1] "THE BOY WHO LIVED Mr. and Mrs. Dursley, of number four, Privet Drive, were proud to say that they were perfectly normal, thank you very much. They were the last people you'd expect to be involved in anything strange or mysterious, because they just didn't hold with such nonsense. Mr. Dursley was the director of a firm called Grunnings, which made drills. He was a big, beefy man with hardly any neck, although he did have a very large mustache. Mrs. Dursley was thin and blonde and had nearly twice the usual amount of neck, which came in very useful as she spent so much of her time craning over garden fences, spying on the neighbors. The Dursleys had a small son called Dudley and in their opinion there was no finer boy anywhere. The Dursleys had everything they wanted, but they also had a secret, and their greatest fear was that somebody would discover it. They didn't think they could bear it if anyone found out about the Potters. Mrs. Potter was Mrs. Dursley's sister, but they h... <truncated>

[2] "THE VANISHING GLASS Nearly ten years had passed since the Dursleys had woken up to find their nephew on the front step, but Privet Drive had hardly changed at all. The sun rose on the same tidy front gardens and lit up the brass number four on the Dursleys' front door; it crept into their living room, which was almost exactly the same as it had

TEXT TO TIBBLE

```
text_tb <- tibble(  
  chapter = seq_along(philosophers_stone),  
  text = philosophers_stone  
)  
text_tb  
# A tibble: 17 x 2  
  chapter text  
  <int> <chr>  
1      1 "THE BOY WHO LIVED    Mr. and Mrs. Dursley, of n...  
2      2 "THE VANISHING GLASS  Nearly ten years had pas...  
3      3 "THE LETTERS FROM NO ONE    The escape of the Br...  
4      4 "THE KEEPER OF THE KEYS    BOOM. They knocked ag...  
5      5 "DIAGON ALLEY    Harry woke early the next morni...  
6      6 "THE JOURNEY FROM PLATFORM NINE AND THREE-QUAR...  
7      7 "THE SORTING HAT    The door swung open at once....  
8      8 "THE POTIONS MASTER    There, look.\"    \"Where?\\...  
9      9 "THE MIDNIGHT DUEL    Harry had never believed h...  
10     10 "HALLOWEEN    Malfoy couldn't believe his eyes w...
```

Creates a tibble with two variables

- chapter
- all text for that chapter

YOUR TURN!

1. Organize `harrypotter::deathly_hallows` into a tibble
2. Organize `harrypotter::chamber_of_secrets` into a tibble
3. Organize `harrypotter::goblet_of_fire` into a tibble

TEXT TO TIBBLE

- But what if we have a series of documents that we want to store and analyze together?

```
df1 <- philosophers_stone  
df2 <- chamber_of_secrets  
df3 <- prisoner_of_azkaban  
df4 <- goblet_of_fire  
df5 <- order_of_the_phoenix  
df6 <- half_blood_prince  
df7 <- deathly_hallows
```

TEXT TO TIBBLE

- But what if we have a series of documents that we want to store and analyze together?

```
df1 <- philosophers_stone
df2 <- chamber_of_secrets
df3 <- prisoner_of_azkaban
df4 <- goblet_of_fire
df5 <- order_of_the_phoenix
df6 <- half_blood_prince
df7 <- deathly_hallows
```

```
ls(pattern = "df")
```

```
[1] "df1" "df2" "df3" "df4" "df5" "df6" "df7"
```

```
books <- mget(ls(pattern = "df"))
```

- **ls** will list all objects with the specified pattern in your global environment
- **mget** will get all the objects specified and hold them in a list

TEXT TO TIBBLE

- Let's set up a process to combine these multiple text objects into a tidied tibble

```
titles <- c(
  "Philosopher's Stone", "Chamber of Secrets",
  "Prisoner of Azkaban", "Goblet of Fire",
  "Order of the Phoenix", "Half-Blood Prince",
  "Deathly Hallows"
)
```

```
series <- tibble()
```

- Step 1: create title names in same order as listed objects
- Step 2: create an empty tibble

TEXT TO TIBBLE

- Let's set up a process to combine these multiple text objects into a tidied tibble

```
titles <- c(
  "Philosopher's Stone", "Chamber of Secrets",
  "Prisoner of Azkaban", "Goblet of Fire",
  "Order of the Phoenix", "Half-Blood Prince",
  "Deathly Hallows"
)
```

```
series <- tibble()
```

```
for(i in seq_along(books)) {
  org <- tibble(
    book      = titles[i],
    chapter   = seq_along(books[[i]]),
    text      = books[[i]]
  )
```

```
  series <- rbind(series, org)
}
```

- Step 1: create title names in same order as listed objects
- Step 2: create an empty tibble
- Step 3: loop through the book list and add the title, chapter, and text for each book and...

...add it to the empty tibble

TEXT TO TIBBLE

- Let's set up a process to combine these multiple text objects into a tidied tibble

```
series
# A tibble: 200 x 3
  book                chapter text
  <chr>              <int> <chr>
1 Philosopher's Stone      1 "THE BOY WHO LIVED    Mr. an...
2 Philosopher's Stone      2 "THE VANISHING GLASS   Near...
3 Philosopher's Stone      3 "THE LETTERS FROM NO ONE ...
4 Philosopher's Stone      4 "THE KEEPER OF THE KEYS B...
5 Philosopher's Stone      5 "DIAGON ALLEY    Harry woke ...
6 Philosopher's Stone      6 "THE JOURNEY FROM PLATFORM...
7 Philosopher's Stone      7 "THE SORTING HAT    The door...
8 Philosopher's Stone      8 "THE POTIONS MASTER   There...
9 Philosopher's Stone      9 "THE MIDNIGHT DUEL    Harry ...
10 Philosopher's Stone     10 "HALLOWEEN    Malfoy couldn'...
# ... with 190 more rows
```

- We end up with a data frame with:
 - Book title
 - Chapter number
 - Chapter text

YOUR TURN!

In your data folder you have three .txt files (amazon_cells_labelled.txt, imdb_labelled.txt, yelp_labelled.txt).

Import these files and get them organized into a data frame that resembles:

file	text
<chr>	<chr>
1 amazon_cells_labelled.txt	So there is no way for me to plug it in here in the US...
2 amazon_cells_labelled.txt	Good case, Excellent value.
3 amazon_cells_labelled.txt	Great for the jawbone.
4 amazon_cells_labelled.txt	Tied to charger for conversations lasting more than 45...
5 amazon_cells_labelled.txt	The mic is great.

UNNESTING

Going from paragraph to words



TEXT TO TIBBLE

```
text_tb <- tibble(  
  chapter = seq_along(philosophers_stone),  
  text = philosophers_stone  
)  
text_tb  
# A tibble: 17 x 2  
  chapter text  
  <int> <chr>  
1      1 "THE BOY WHO LIVED    Mr. and Mrs. Dursley, of n...  
2      2 "THE VANISHING GLASS  Nearly ten years had pas...  
3      3 "THE LETTERS FROM NO ONE    The escape of the Br...  
4      4 "THE KEEPER OF THE KEYS    BOOM. They knocked ag...  
5      5 "DIAGON ALLEY    Harry woke early the next morni...  
6      6 "THE JOURNEY FROM PLATFORM NINE AND THREE-QUAR...  
7      7 "THE SORTING HAT    The door swung open at once....  
8      8 "THE POTIONS MASTER    There, look.\"    \"Where?\\...  
9      9 "THE MIDNIGHT DUEL    Harry had never believed h...  
10     10 "HALLOWEEN    Malfoy couldn't believe his eyes w...
```

We now have our book in a data frame with each chapter text in its own row.

UNNEST OUR TEXT

```
text_tb %>% unnest_tokens(word, text)
text_tb %>% unnest_tokens(word, text, token = "sentences")
text_tb %>% unnest_tokens(word, text, token = "lines")
text_tb %>% unnest_tokens(word, text, token = "ngrams", n = 2)
```

```
# A tibble: 77,858 × 2
```

	chapter	word
	<int>	<chr>
1	1	the boy
2	1	boy who
3	1	who lived
4	1	lived mr
5	1	mr and
6	1	and mrs
7	1	mrs duncey

We can unnest our text by:

- individual words (default)
- sentences (looks for ".")
- lines
- ngrams

Unnest will also:

- standardize to lowercase
- remove non-alphanumeric chars

YOUR TURN!

*Using the **deathly_hallows** data, unnest the text by single words and then bi-grams*

STOP WORDS

Where's the content?

STOP WORDS

Word	Word	Word	Word	Word	Word
i	by	they're	his	then	a
yours	after	they'd	them	both	while
herself	off	weren't	these	not	through
which	where	shouldn't	have	ourselves	in
was	other	there's	could	she	here
does	than	if	you've	theirs	few
she's	my	with	you'll	am	own
he'd	yourselves	below	haven't	had	your
isn't	its	under	cannot	i'm	hers
won't	whom	how	where's	they've	what
who's	be	such	because	she'll	are
the	doing	very	against	doesn't	do
at	we're	we	from	mustn't	he's
before	we'd	him	further	how's	you'd
on	wasn't	they	any	until	they'll
when	shan't	that	nor	into	didn't
most	here's	being	ours	down	that's
so	but	should	himself	once	an

- Stop words are common english words
- Typically do not provide us much context

STOP WORDS

```
tidytext::stop_words
# A tibble: 1,149 × 2
  word lexicon
  <chr>   <chr>
1      a    SMART
2     a's    SMART
3    able    SMART
4   about    SMART
5   above    SMART
6 according    SMART
7 accordingly    SMART
8   across    SMART
9  actually    SMART
10   after    SMART
# with 1,139 more rows
```

- Stop words are common english words
- Typically do not provide us much context
- **tidytext** provides stop words from three lexicons
 - onix
 - SMART
 - snowball

REMOVING STOP WORDS

```
text_tb %>%  
  unnest_tokens(word, text) %>%  
  count(word, sort = TRUE)  
# A tibble: 5,978 x 2  
  word      n  
  <chr> <int>  
1 the     3629  
2 and     1923  
3 to      1861  
4 a       1691  
5 he      1528  
6 of      1267  
7 harry   1213  
8 was     1186  
9 it      1037
```

- If we do not remove stop words, any frequency analysis will be dominated by non-contextual words.

REMOVING STOP WORDS

```
text_tb %>%  
  unnest_tokens(word, text) %>%  
  anti_join(stop_words) %>%  
  count(word, sort = TRUE)  
# A tibble: 5,421 × 2  
  word      n  
  <chr> <int>  
1  harry 1213  
2   ron  410  
3 hagrid 336  
4 hermione 257  
5 professor 181  
6   looked 169  
7   snape 145  
8 dumbledore 143
```

- If we did not remove stop words, any frequency analysis will be dominated by non-contextual words.
- **anti_join** returns all words in our text that are not in the **stop_words** list

YOUR TURN!

*Unnest the **deathly_hallows** book and remove the stop words.*

FREQUENCIES

What are the most common words?



WORD FREQUENCY

```
text_tb %>%  
  unnest_tokens(word, text) %>%  
  anti_join(stop_words) %>%  
  count(word, sort = TRUE) %>%  
  top_n(10)
```

```
# A tibble: 10 x 2
```


	word	n
	<chr>	<int>
1	harry	1213
2	ron	410
3	hagrid	336
4	hermione	257
5	professor	181
6	looked	169
7	space	145

- **count** counts all occurrences of each word (returns **n**).
- **top_n** returns the top **n** words

WORD FREQUENCY

```
text_tb %>%  
  unnest_tokens(word, text) %>%  
  anti_join(stop_words) %>%  
  count(word, sort = TRUE) %>%  
  top_n(10) %>%  
  ggplot(aes(reorder(word, n), n)) +  
  geom_col() +  
  coord_flip()
```

We can pipe this right into ggplot



N-GRAM FREQUENCY

```
text_tb %>%
  unnest_tokens(bigram, text, token = "ngrams", n = 2) %>%
  count(bigram, sort = TRUE)
# A tibble: 42,756 x 2
   bigram      n
   <chr> <int>
1  of the    285
2  in the    269
3  on the    218
4  it was    207
5  he was    194
6  to the    173
7  out of    148
8  at the    141
9  said hanny 136
```

- We can use the same process to get bi-gram frequencies...
- ...but we have a stop word problem again.

How can we remove
these?

N-GRAM FREQUENCY

```
text_tb %>%  
  unnest_tokens(bigram, text, token = "ngrams", n = 2) %>%  
  separate(bigram, c("word1", "word2"), sep = " ")
```

```
# A tibble: 77,858 x 3
```

	chapter	word1	word2
*	<int>	<chr>	<chr>
1	1	the	boy
2	1	boy	who
3	1	who	lived
4	1	lived	mr
5	1	mr	and
6	1	and	mrs
7	1	mrs	dursley
8	1	dursley	of
9	1	of	number

- First, we use **separate** (tidyr package) to separate the bi-gram.

N-GRAM FREQUENCY


```
text_tb %>%
  unnest_tokens(bigram, text, token = "ngrams", n = 2) %>%
  separate(bigram, c("word1", "word2"), sep = " ") %>%
  filter(
    !word1 %in% stop_words$word,
    !word2 %in% stop_words$word
  ) %>%
  count(word1, word2, sort = TRUE)
# A tibble: 7,695 x 3
   word1      word2      n
   <chr>    <chr> <int>
1   uncle    vernon    97
2 professor mcgonagall  90
3    aunt   petunia    52
4   harry    potter    26
```

- First, we use **separate** (tidyr package) to separate the bi-gram.
- Then we can **filter** out stop words from each word variable and
- **count** by word1 & word2 pairs.

N-GRAM FREQUENCY

```
text_tb %>%  
  unnest_tokens(bigram, text, token = "ngrams", n = 2) %>%  
  separate(bigram, c("word1", "word2"), sep = " ") %>%  
  filter(  
    !word1 %in% stop_words$word,  
    !word2 %in% stop_words$word  
  ) %>%  
  count(word1, word2, sort = TRUE) %>%  
  unite(bigram, word1, word2, sep = " ") %>%  
  top_n(10) %>%  
  ggplot(aes(reorder(bigram, n), n)) +  
  geom_col() +  
  coord_flip()
```

We can pipe this right into ggplot



YOUR TURN!

*Find the most common tri-grams in **deathly_hallows**
(extra credit for those that can visualize the output).*

COMPARISONS

Comparing frequencies across multiple documents



HIGHLY CORRELATED WORDS WITHIN A DOCUMENT

```
ph_cor <- tibble(  
  title = "Philosopher's Stone",  
  chapter = seq_along(philosophers_stone),  
  text = philosophers_stone  
) %>%  
  unnest_tokens(word, text) %>%  
  anti_join(stop_words) %>%  
  filter(n() >= 50) %>%  
  widyr::pairwise_cor(word, chapter, sort = TRUE)
```

```
filter(ph_cor, item1 == "harry")
```

```
# A tibble: 5,420 x 3
```

	item1	item2	correlation
	<chr>	<chr>	<dbl>
1	harry	gloom	0.761
2	harry	harry's	0.746
3	harry	paper	0.734

- **pairwise_cor** assess words that are most correlated with other words.

COMPARING FREQUENCIES ACROSS BOOKS

```
series %>%  
  unnest_tokens(word, text) %>%  
  anti_join(stop_words) %>%  
  count(book, word, sort = TRUE)
```

```
# A tibble: 63,651 x 3
```

	book	word	n
	<chr>	<chr>	<int>
1	Order of the Phoenix	harry	3730
2	Goblet of Fire	harry	2936
3	Deathly Hallows	harry	2770
4	Half-Blood Prince	harry	2581
5	Prisoner of Azkaban	harry	1824
6	Chamber of Secrets	harry	1503
7	Order of the Phoenix	hermione	1220
8	Philosopher's Stone	harry	1212

- What about comparing frequencies across multiple documents?
- What issue may we run into?

COMPARING FREQUENCIES ACROSS BOOKS

```
books %>%  
  map(str_count) %>%  
  map_dbl(sum)  
  df1      df2      df3      df4      df5      df6      df7      dh_df  
439430 490784 611209 1105151 1514903 984113 1124196 1137387
```

- Total word counts across our documents range from 439,430 to 1,514,903
- Longer books will dominate any word frequency comparisons.
- We can use proportions to compare.

COMPARING TWO DOCUMENTS

```
df1 <- tibble(  
  title   = "Philosopher's Stone",  
  chapter = seq_along(philosophers_stone),  
  text    = philosophers_stone  
) %>%  
  unnest_tokens(word, text) %>%  
  anti_join(stop_words)
```

```
df2 <- tibble(  
  title   = "Deathly Hallows",  
  chapter = seq_along(deathly_hallows),  
  text    = deathly_hallows  
) %>%  
  unnest_tokens(word, text) %>%  
  anti_join(stop_words)
```

- Step 1: create unnested, clean data for each document

COMPARING TWO DOCUMENTS

```
combined <- df1 %>%
  rbind(df2) %>%
  filter(str_detect(word, "[a-z']+")) %>%
  count(title, word) %>%
  group_by(title) %>%
  mutate(pct = n / sum(n)) %>%
  select(-n) %>%
  spread(title, pct) %>%
  na.omit() %>%
  mutate(delta = abs(`Deathly Hallows` - `Philosopher's Stone`))
```

```
combined
# A tibble: 3,880 x 4
  word      `Deathly Hallows` `Philosopher's Stone` delta
  <chr>          <dbl>          <dbl>    <dbl>
1 aargh          0.0000136          0.0000350 0.0000214
2 aback          0.0000681          0.0000350 0.0000331
3 abbott         0.0000681          0.0000350 0.0000331
4 abnormal       0.0000136          0.0000350 0.0000214
5 absolutely     0.0000818          0.0000700 0.0000117
6 absurd         0.0000136          0.0000350 0.0000214
7 accept         0.000177           0.0000700 0.000107
8 accepted       0.0000818          0.0000700 0.0000117
9 accident       0.000123           0.0000350 0.0000876
10 accidentally  0.0000409          0.0000700 0.0000291
# ... with 3.870 more rows
```


- Step 1: create unnested, clean data for each document
- Step 2:
 - **rbind** data frames together
 - **Filter out non-words**
 - **Count words and compute proportion for each book**
 - **Spread data**

COMPARING TWO DOCUMENTS

```
# we can plot
ggplot(combined, aes(x = `Philosopher's Stone`, y = `Deathly Hallows`, color = delta)) +
  geom_abline(color = "gray40", lty = 2) +
  geom_jitter(alpha = 0.1, size = 2.5, width = 0.2, height = 0.3) +
  geom_text(aes(label = word), check_overlap = TRUE) +
  scale_x_log10("Philosopher's Stone", labels = scales::percent) +
  scale_y_log10("Deathly Hallows", labels = scales::percent) +
  scale_color_gradient(limits = c(0, 0.001), low = "darkslategray4", high = "gray75") +
  theme(legend.position = "none")
```

```
# or compute correlation
cor.test(combined$`Deathly Hallows`, combined$`Philosopher's Stone`)
Pearson's product-moment correlation
```

```
data: combined$`Deathly Hallows` and combined$`Philosopher's Stone`
t = 132.6, df = 3878, p-value < 2.2e-16
alternative hypothesis: true correlation is not equal to 0
95 percent confidence interval:
 0.8992951 0.9106770
sample estimates:
      cor
0.9051481
```



COMPARING 3+ DOCUMENTS

What about comparing across three or more?

COMPARING 3+ DOCUMENTS

```
# 1. clean
clean_tokens <- series %>%
  unnest_tokens(word, text) %>%
  anti_join(stop_words)

clean_tokens
# A tibble: 409,338 x 3
   book                chapter word
   <chr>              <int> <chr>
1 Philosopher's Stone      1 boy
2 Philosopher's Stone      1 lived
3 Philosopher's Stone      1 dursley
4 Philosopher's Stone      1 privet
5 Philosopher's Stone      1 drive
6 Philosopher's Stone      1 proud
7 Philosopher's Stone      1 perfectly
8 Philosopher's Stone      1 normal
9 Philosopher's Stone      1 people
10 Philosopher's Stone     1 expect
# ... with 409,328 more rows
```

- Step I: Unnest and remove stop words.

COMPARING 3+ DOCUMENTS

```
# 1. clean
clean_tokens <- series %>%
  unnest_tokens(word, text) %>%
  anti_join(stop_words)

# 2. compute percent of word use across all novels
series_pct <- clean_tokens %>%
  count(word, sort = TRUE) %>%
  transmute(word, all_words = n / sum(n))
```

```
series_pct
# A tibble: 23,795 x 2
   word      all_words
  <chr>      <dbl>
1 harry      0.0404
2 ron        0.0140
3 hermione    0.0120
4 dumbledore 0.00702
5 looked     0.00573
6 professor  0.00490
7 hagrid     0.00423
8 time       0.00418
9 wand       0.00400
10 eyes      0.00392
# ... with 23,785 more rows
```

- Step 1: Unnest and remove stop words.
- Step 2: Compute % of word use across series

COMPARING 3+ DOCUMENTS

```
# 1. clean
clean_tokens <- series %>%
  unnest_tokens(word, text) %>%
  anti_join(stop_words)

# 2. compute percent of word use across all novels
series_pct <- clean_tokens %>%
  count(word, sort = TRUE) %>%
  transmute(word, all_words = n / sum(n))

# 3. compute percent of word use within each novel
comp_prop <- clean_tokens %>%
  count(book, word, sort = TRUE) %>%
  group_by(book) %>%
  mutate(book_words = n / sum(n)) %>%
  inner_join(series_pct) %>%
  ungroup()

comp_prop
# A tibble: 63,651 x 5
   book          word      n book_words all_words
  <chr>        <chr> <int>      <dbl>    <dbl>
1 Order of the Phoenix harry   3730    0.0385    0.0404
2 Goblet of Fire      harry   2936    0.0404    0.0404
3 Deathly Hallows     harry   2770    0.0377    0.0404
4 Half-Blood Prince   harry   2581    0.0409    0.0404
```

- Step 1: Unnest and remove stop words.
- Step 2: Compute % of word use across series
- Step 3: Compute % of word use within each novel & join series %.

Now what can we do?

COMPARING 3+ DOCUMENTS

```
# option 1: Plot
comp_prop %>%
  mutate(
    delta = abs(all_words - book_words),
    book = factor(book, levels = titles)
  ) %>%
  group_by(book) %>%
  top_n(100, wt = delta) %>%
  ggplot(aes(x = all_words, y = book_words, color = delta)) +
  geom_abline(color = "gray40", lty = 2) +
  geom_jitter(alpha = 0.1, size = 2.5, width = 0.3, height = 0.3) +
  geom_text(aes(label = word), check_overlap = TRUE) +
  scale_x_log10("Proportion of series", labels = scales::percent) +
  scale_y_log10("Proportion of book", labels = scales::percent) +
  scale_color_gradient(limits = c(0, 0.001), low = "darkslategray4", high = "gray75") +
  facet_wrap(~ book, ncol = 2) +
  theme(legend.position = "none")
```



COMPARING 3+ DOCUMENTS

```
# option 2: Compute correlation of each book to the overall series
```

```
comp_prop %>%
```

```
  group_by(book) %>%
```

```
  summarise(rho = cor(book_words, all_words)) %>%
```

```
  arrange(desc(rho))
```

```
# A tibble: 7 x 2
```

```
  book          rho
```

```
  <chr>        <dbl>
```

```
1 Order of the Phoenix 0.984
```

```
2 Goblet of Fire      0.979
```

```
3 Deathly Hallows    0.970
```

```
4 Half-Blood Prince  0.970
```

```
5 Chamber of Secrets 0.966
```

```
6 Prisoner of Azkaban 0.964
```


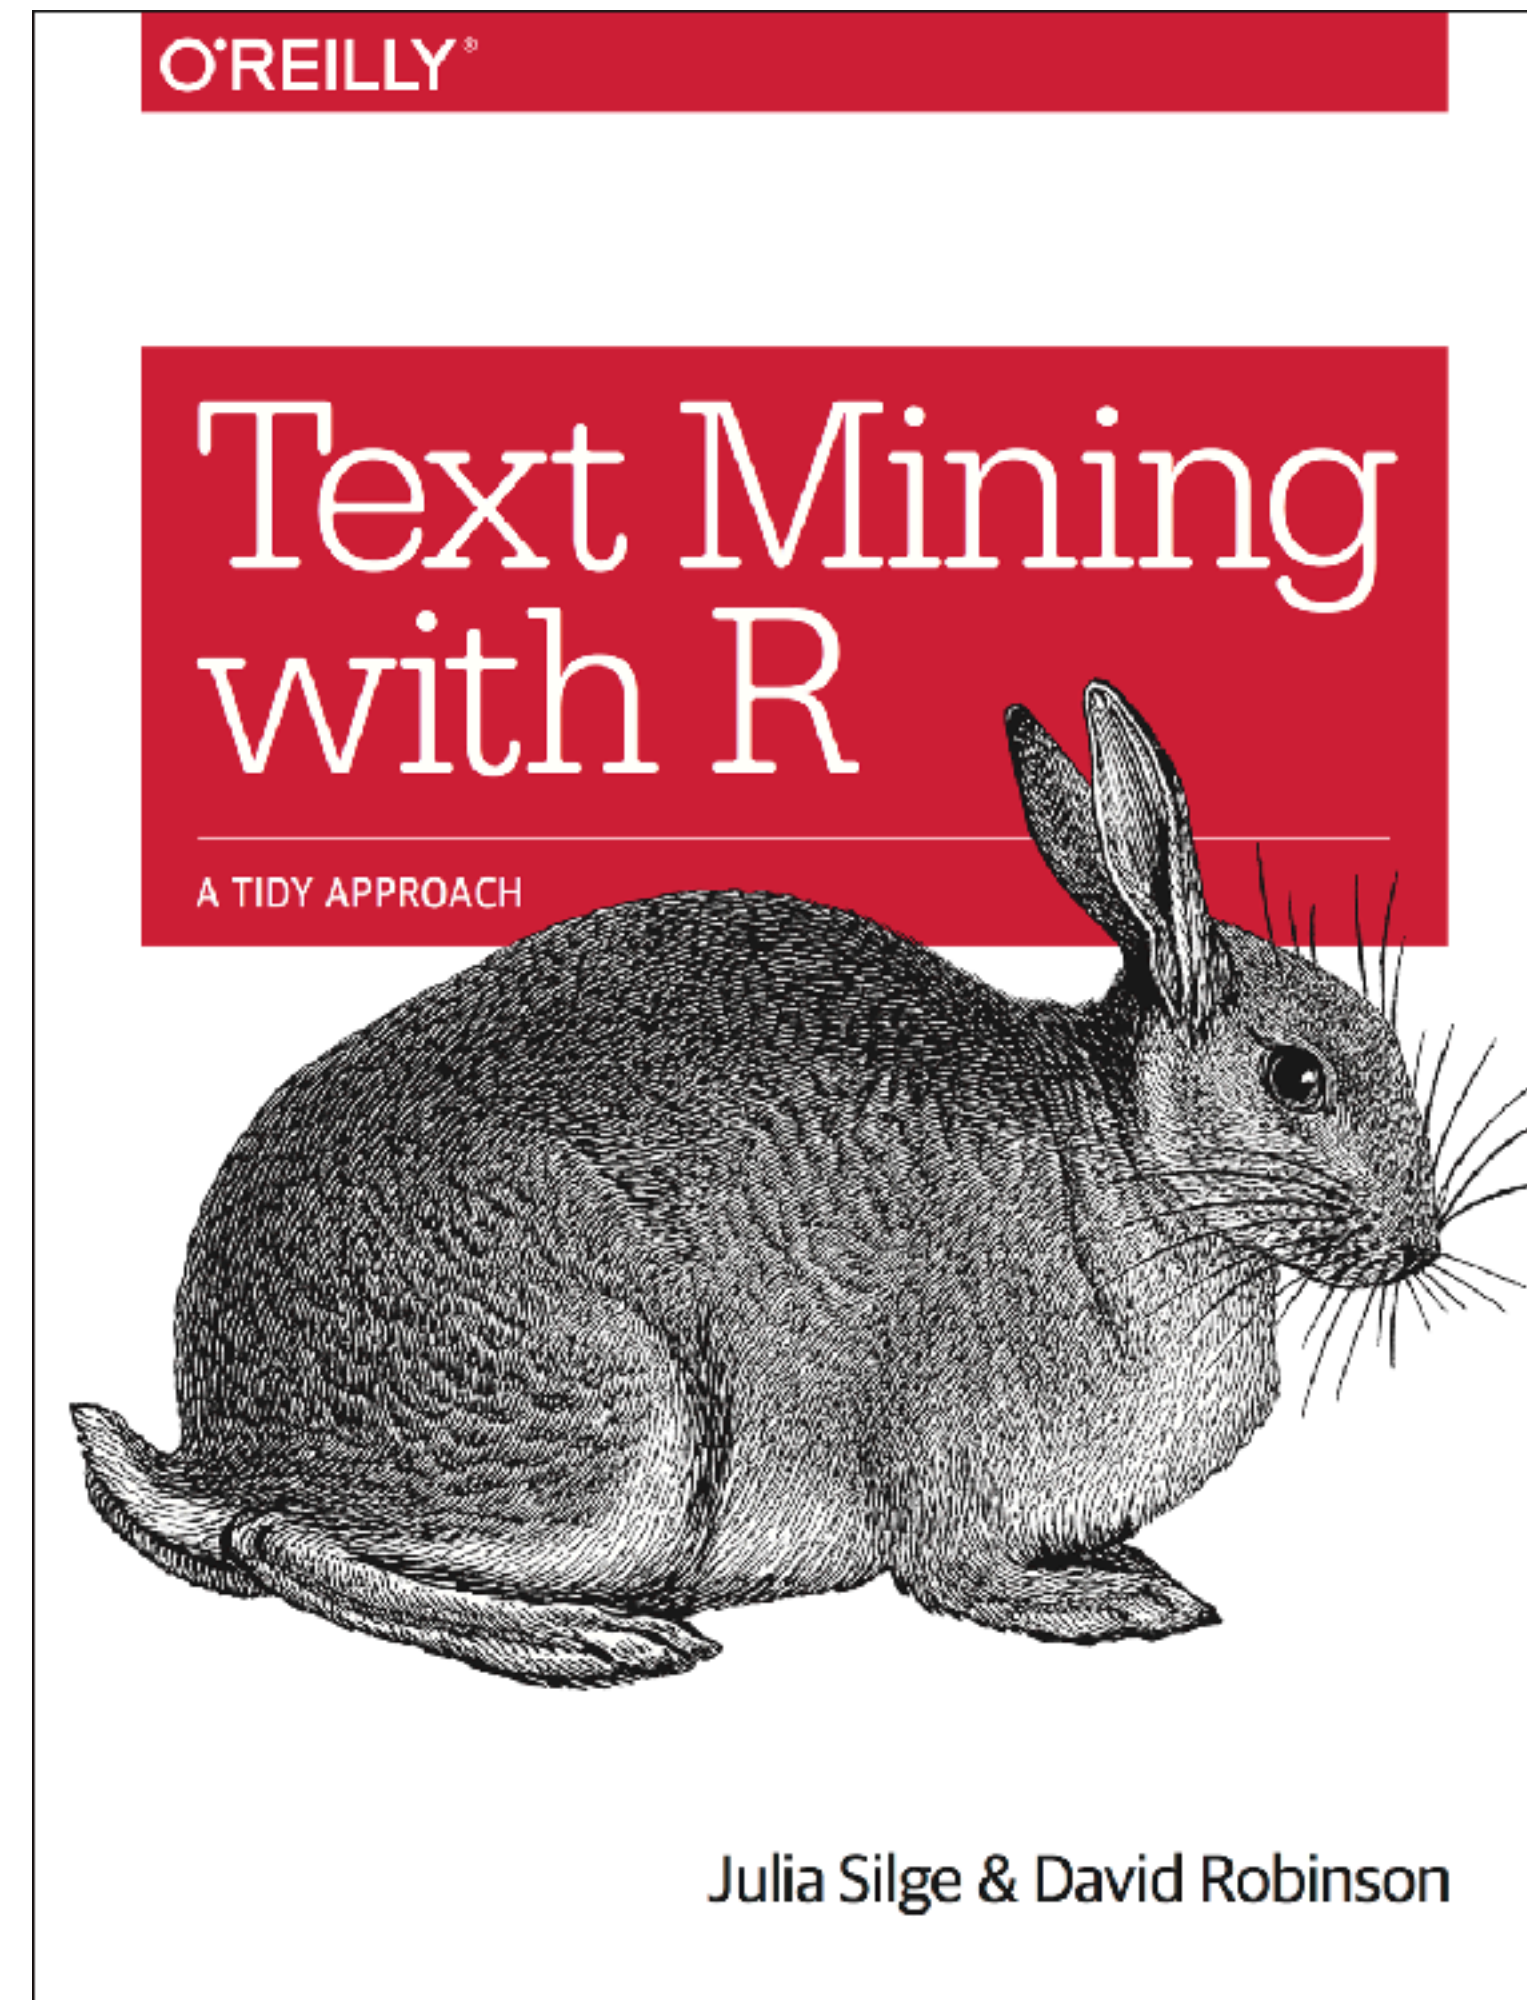

```
7 Philosopher's Stone 0.955
```

- Philosopher's stone is least representative of the overall series.

SO LITTLE TIME!



LEARN MORE



WHAT TO REMEMBER



FUNCTIONS TO REMEMBER

Operator/Function	Description
<code>tibble</code>	Convert character strings into a data frame (tibble)
<code>unnest_tokens</code>	Unnest and clean text
<code>stop_words</code>	Commonly used English words that lack context
<code>anti_join</code> , <code>filter</code> , <code>count</code>	dplyr functions commonly used