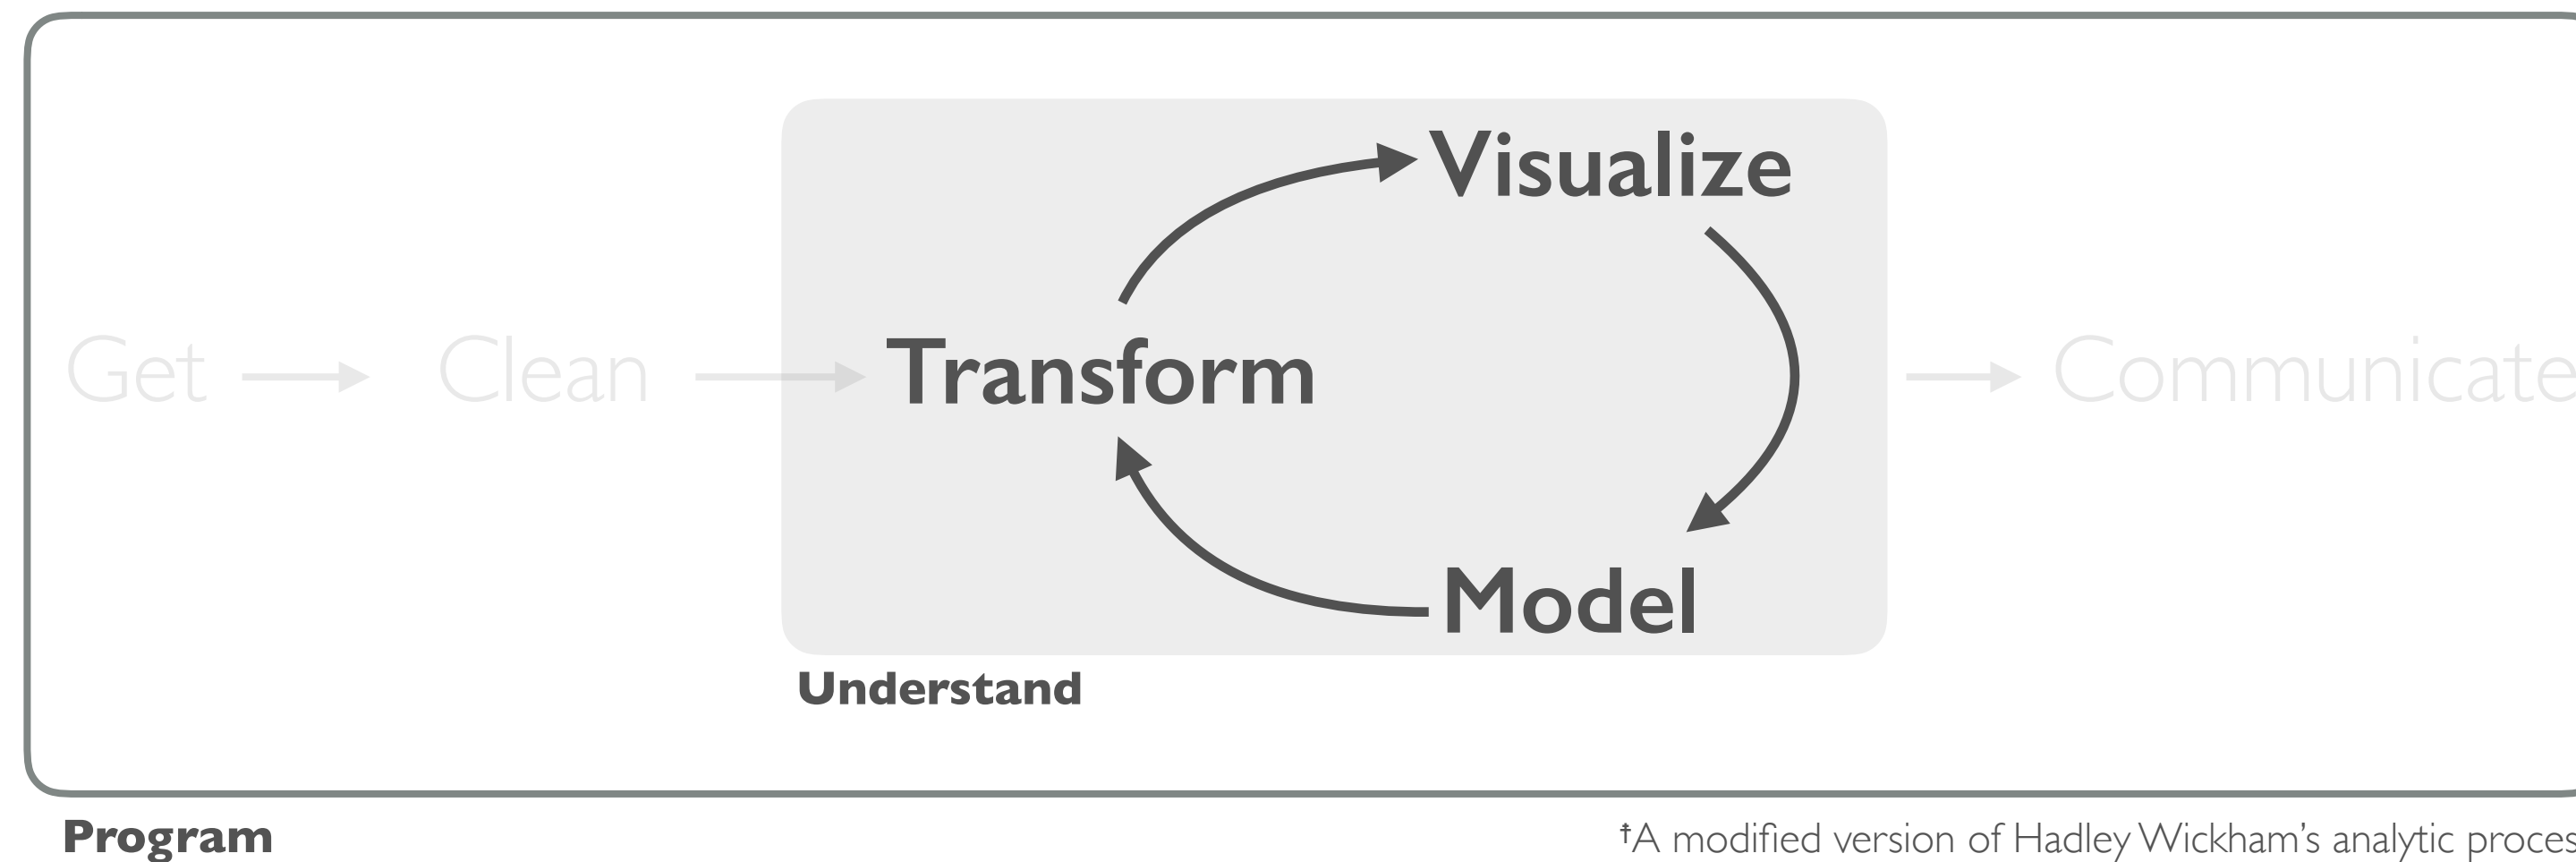



PREDICTIVE MODELING






What society thinks I do



What my friends think I do



What other computer scientists think I do



What mathematicians think I do



What I think I do


```
Source on Save  
# Load required pkgs  
library(randomForest)  
library(xgboost)
```

What I actually do

PREDICTIVE MODELING

- Text data can be used for predictive modeling much like more traditional data.
- Can be used for both classification and regression problems
- Caveats
 - Must be ***cleaned*** and ***tidied*** (document term matrix)
 - Technique must be able to handle the ***curse of dimensionality***
 - Technique must be able to handle ***sparsity***

PREREQUISITES



PACKAGE PREREQUISITE

```
library(tidyverse)      # data wrangling & plotting
library(tidytext)       # text manipulation
library(glmnet)         # elastic net modeling
library(pROC)           # ROC curve / AUC
```

DATA PREREQUISITE

```
# Click-bait news headlines
```

```
url <- "https://raw.githubusercontent.com/kwartler/text_mining/master/all_3k_headlines.csv"
```

```
headlines <- read_csv(url)
```

```
headlines
```

```
# A tibble: 3,000 x 4
```

	headline	url	site	y
	<chr>	<chr>	<chr>	<int>
1	Mom Sentenced To 6 Years In Prison Fo...	http://dailybuzzlive.com/mom-sentenced-to...	dailyb...	1
2	"The Most Shocking '\x98Jerry Springe...	http://dailybuzzlive.com/the-most-shockin...	dailyb...	1
3	America''s Self-Inflicted Defense Woes	http://www.activistpost.com/2016/08/ameri...	activi...	1
4	"A Man Spots A Reckless Driver, When ...	http://dailybuzzlive.com/a-man-spots-a-re...	dailyb...	1
5	Tim Cook Asks FBI To Withdraw Order T...	http://www.buzzfeed.com/johnpaczkowski/ap...	buzzfe...	1
6	Police Union Threatens Not To Patrol ...	http://www.buzzfeed.com/salvadorhernandez...	buzzfe...	1
7	Ed Rendell: Clinton's Speech Was Pres...	http://www.buzzfeed.com/christophermassie...	buzzfe...	1

DATA PREP

Preparing for the modeling process



GENERALIZABILITY

- **Generalizability**: the confidence that a model will perform similarly on unobserved data.
- The difference between a descriptive vs. predictive capability


A more generalizable model gives us greater confidence in future accuracy

BIAS vs. VARIANCE

Bias

- Wrong model assumptions
- Model restrictions on hyperparameters
- Error due to model specification

Underfitting




BIAS _{vs.} VARIANCE

Bias

- Wrong model assumptions
- Model restrictions on hyperparameters
- Error due to model specification


Underfitting



Variance

- Model over adapts to training data
- Hyperparameters overly tuned to training data
- Error due to the sampling of training set

Overfitting



BIAS _{vs.} VARIANCE


Bias

- Wrong model assumptions
- Model restrictions on hyperparameters
- Error due to model specification


Variance

- Model over adapts to training data
- Hyperparameters overly tuned to training data
- Error due to the sampling of training set


Underfitting



Bias Variance
Tradeoff



Overfitting




BIAS vs. VARIANCE


Resampling

- The more adaptable our modeling techniques become, the greater the importance to use advanced resampling methods to minimize overfitting... **aka over-confident predictions**


Underfitting



Bias Variance
Tradeoff




Overfitting



BASIC TRAIN/TEST SPLIT

- At a minimum, always use a training-test split to assess out-of-sample accuracy.
- Typical: 70/30
 - Dependent on size of data set



Disadvantages

- Model can overfit to training data
- Single test set provides only a single insight into predictive performance

ESTABLISHING TRAINING VS. TESTING SETS

```
# assign about 70% to training and 30% to test  
set.seed(123)
```

```
headlines <- headlines %>%  
  mutate(partition = sample(  
    x = c("Train", "Test"),  
    size = 3000,  
    prob = c(.7, .3),  
    replace = TRUE)  
  )
```

`set.seed` provides reproducibility

Next, create a `partition` variable

- sample from `x`
- 3000 times (length of data set)
- 70% from Train, 30% from Test

CLEAN AND TIDY DATA

```
# assign about 70% to training and 30% to test
headlines_tidy <- headlines %>%
  unnest_tokens(word, headline) %>%
  anti_join(stop_words) %>%
  count(url, partition, y, word) %>%
  mutate(id = paste(url, partition, sep = "_"))
```

headlines_tidy

A tibble: 18,647 x 6

	url	partition	y	word	n	id
	<chr>	<chr>	<int>	<chr>	<int>	<chr>
1	http://21stce...	Train	1	anti	1	http://21stce...
2	http://21stce...	Train	1	gove...	1	http://21stce...
3	http://21stce...	Train	1	inst...	1	http://21stce...
4	http://21stce...	Train	1	lead...	1	http://21stce...
5	http://21stce...	Train	1	obama	1	http://21stce...
6	http://21stce...	Train	1	opin...	1	http://21stce...
7	http://21stce...	Train	1	party	1	http://21stce...

This is performing the same tidying process you've seen.

However, we add one step where we **combine the url & partition**:

id: **http://21stcenturywire..._Train**

We will use this later!

CONVERT TO A DOCUMENT TERM MATRIX

```
# convert to a DTM matrix
headlines_matrix <- headlines_tidy %>%
  cast_dtm(id, word, n) %>%
  as.matrix()
```

To model we need to have our data in a DTM style but as a matrix object.

Now we're ready to create our training & testing splits

TRAINING SPLIT

```
train.x <- headlines_matrix %>%  
  as_data_frame() %>%  
  mutate(id = row.names(headlines_matrix)) %>%  
  filter(str_detect(id, "_Train")) %>%  
  select(-id) %>%  
  as.matrix() %>%  
  Matrix::Matrix(sparse = TRUE)
```

```
train.y <- headlines_tidy %>%  
  filter(partition == "Train") %>%  
  distinct(id, y) %>%  
  . $y
```

```
train.x[1:5, 1:5]  
5 x 5 sparse Matrix of class "dgCMatrix"  
      anti government instructs leader obama  
[1,]    1          1          1          1          1  
[2,]    .          .          .          .          .  
[3,]    .          .          .          .          .
```

Create training feature set, which is in the form of a [sparse matrix](#)... computationally efficient!

Create response variable vector

YOUR TURN!

You created the training split, now can you do the same thing to create the test split?

MODELING





Regularized regression: Ridge, Lasso, & Elastic Nets



LINEAR REGRESSION

Goal of linear regression is to fit a line to the data that minimizes total squared error

Minimize RSS:
$$\text{RSS} = \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2$$



LINEAR REGRESSION

Goal of linear regression is to fit a line to the data that minimizes total squared error

- Assumptions:
 - Linear relationship
 - Multivariate normality
 - No or little multicollinearity
 - No autocorrelation
 - Homoscedastic (constant variance in residuals)
 - $p < n$ (there is no unique solution when $p > n$)

LINEAR REGRESSION

Goal of linear regression is to fit a line to the data that minimizes total squared error

- Assumptions:

- Linear relationship
- Multivariate normality
- No or little multicollinearity ✓
- No autocorrelation
- Homoscedastic (constant variance in residuals)
- $p < n$ (there is no unique solution when $p > n$) ✓

*Two problems we run into with text
Both cause serious problems with model stability
& reliability*

REGULARIZED REGRESSION

- Three **shrinkage methods** have been developed to resolve these (and other) concerns: Lasso, Ridge regression, Elastic nets.
- **Constrains** or **regularizes** the coefficient estimates, or equivalently, that **shrinks** the coefficient estimates towards zero.
- It may not be immediately obvious why such a constraint should improve the fit, but it turns out that shrinking the coefficient estimates can significantly reduce their variance.

REGULARIZED REGRESSION

- Three **shrinkage methods** have been developed to resolve these (and other) concerns: Lasso, Ridge regression, Elastic nets.
- **Constrains** or **regularizes** the coefficient estimates, or equivalently, that **shrinks** the coefficient estimates towards zero.

- Lasso minimizes:
$$\text{RSS} + \lambda \sum_{j=1}^p |\beta_j|.$$


- Ridge minimizes:
$$\text{RSS} + \lambda \sum_{j=1}^p \beta_j^2.$$

lambda (λ) becomes a tuning parameter:


- $\lambda \approx 0$: equates to OLS regression
- $\lambda \approx \infty$: equates to the mean as all coefficients are forced to 0

REGULARIZED REGRESSION

Ridge Regression

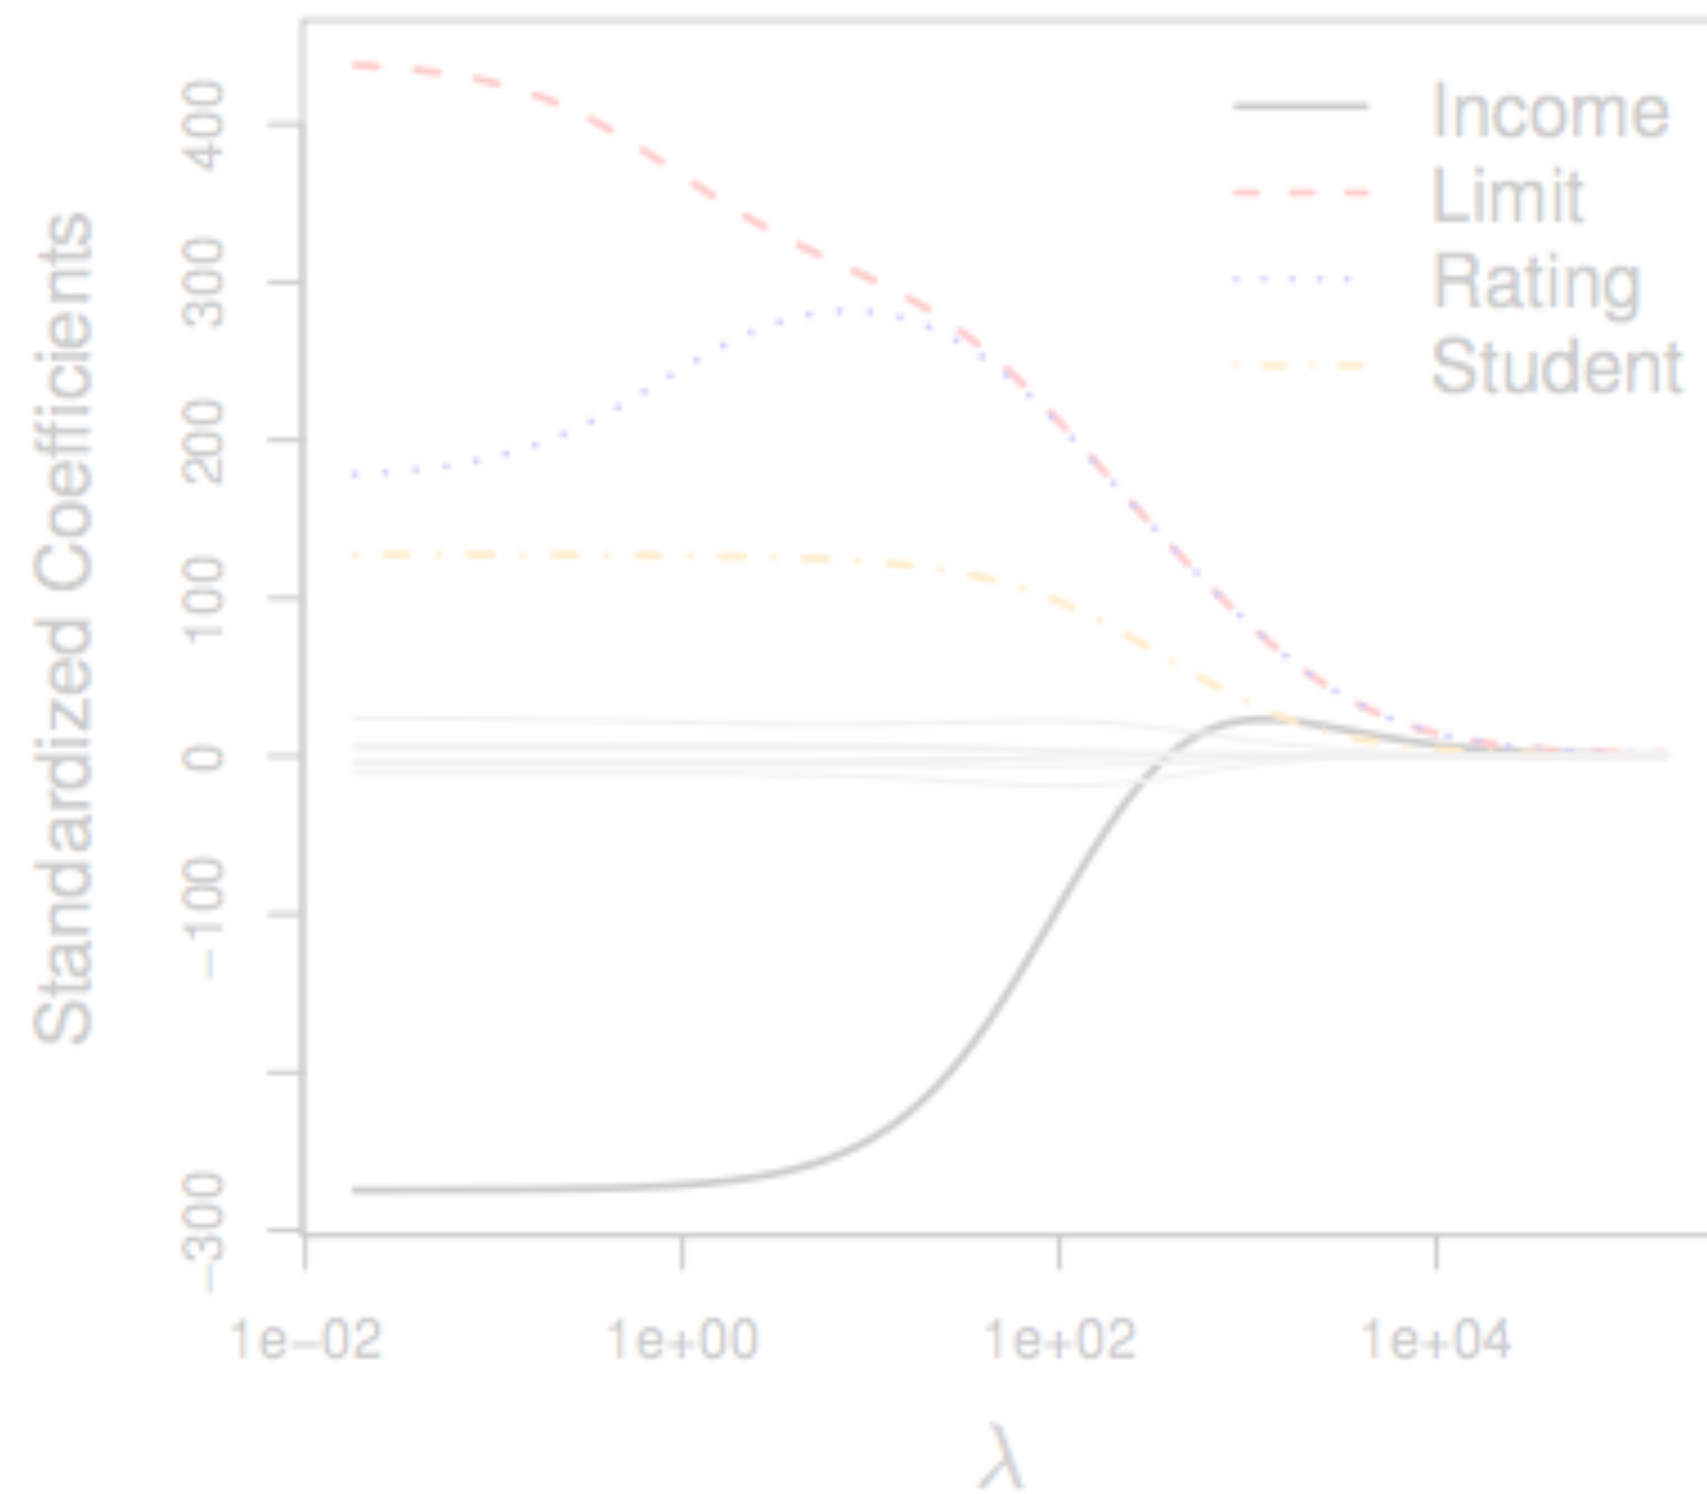


Lasso Regression



REGULARIZED REGRESSION


Ridge Regression



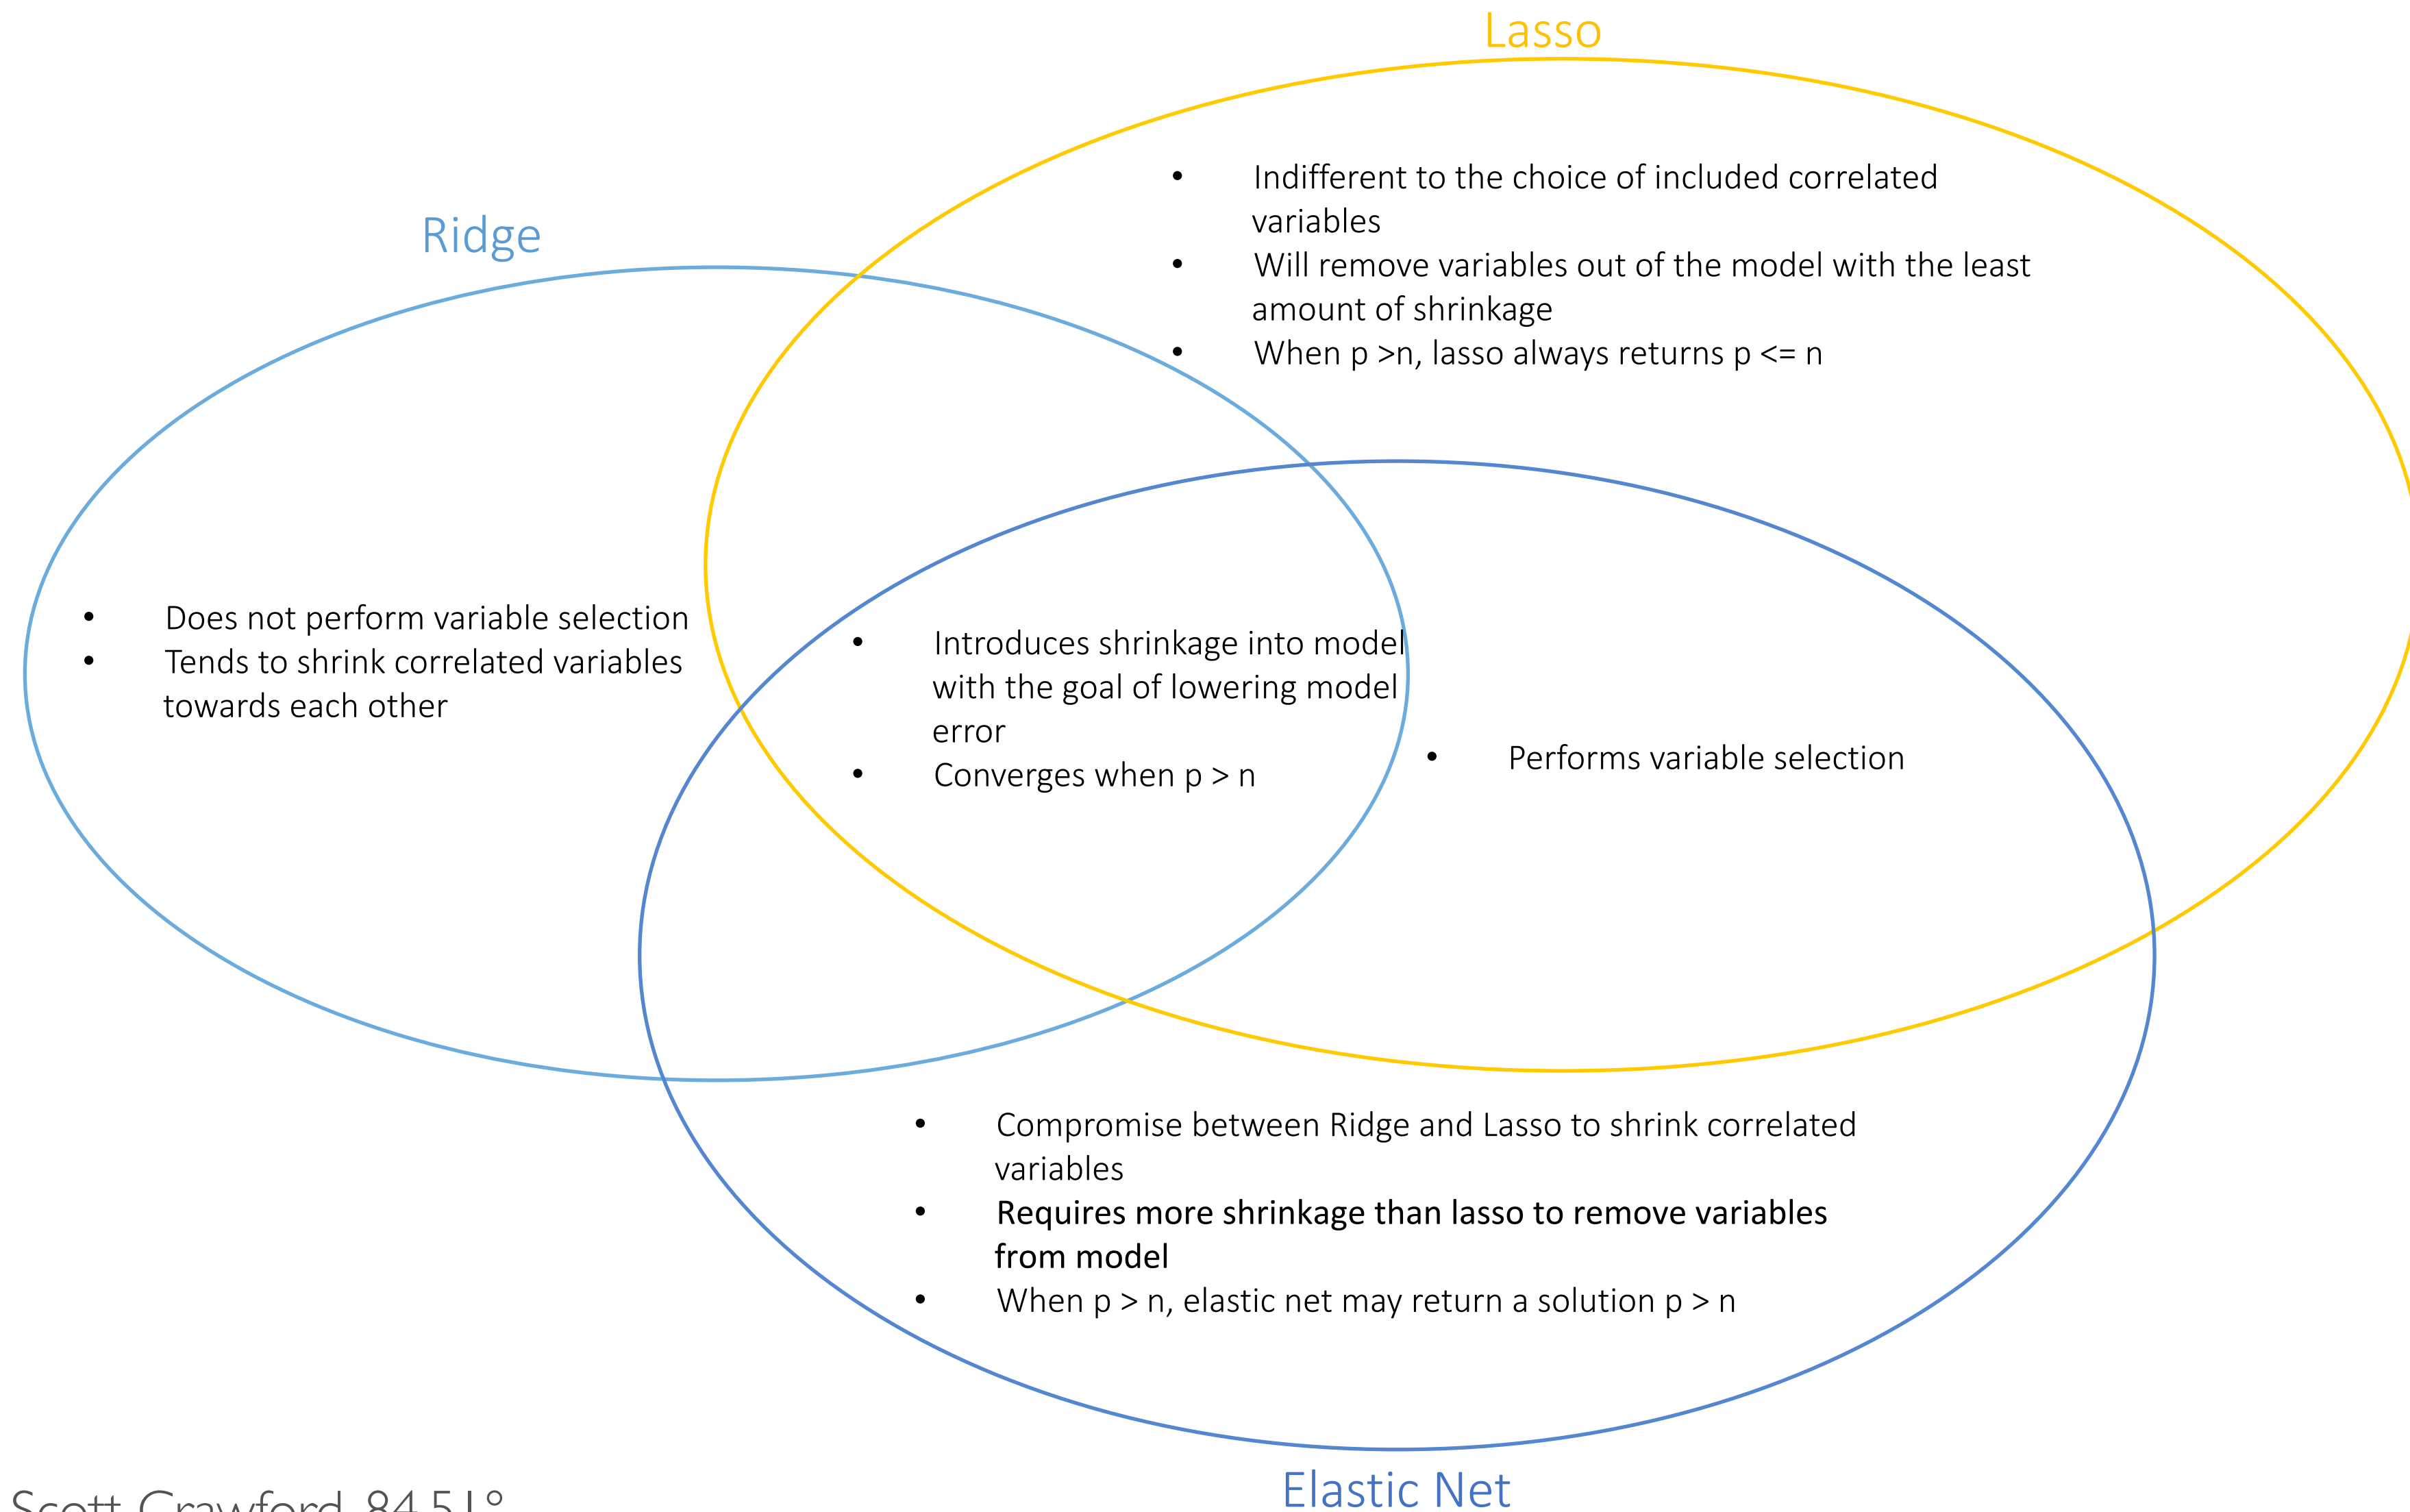
Elastic Net



Lasso Regression



REGULARIZED REGRESSION



REGULARIZED REGRESSION

Q: When do I use one versus the other?

When to use Ridge:

- You have reason to believe all coefficients should be kept in the model.
- You want correlated variables to be shrunk towards each other.
- Lower error rate than Lasso and Elastic Net

When to use Lasso:

- Perform variable selection - You have reason to believe a small subset of variables should be included in the final model
- $p \gg n$
- Highly correlated variables that can be “arbitrarily” dropped
- Lower error rate than Ridge and Elastic Net

When to use Elastic net:

- A balance between the effects of Ridge and Lasso
- $p \gg n$ but okay with a solution $p > n$
- Lower error rate than Ridge and Lasso
- This is currently the most popular version of regularized regression

REGULARIZED REGRESSION

Q: When do I use one versus the other?

A: Most of the time try them all!

When to use Ridge:

- You have reason to believe all coefficients should be kept in the model.
- You want correlated variables to be shrunk towards each other.
- Lower error rate than Lasso and Elastic Net

When to use Lasso:

- Perform variable selection - You have reason to believe a small subset of variables should be included in the final model
- $p \gg n$
- Highly correlated variables that can be “arbitrarily” dropped
- Lower error rate than Ridge and Elastic Net

When to use Elastic net:

- A balance between the effects of Ridge and Lasso
- $p \gg n$ but okay with a solution $p > n$
- Lower error rate than Ridge and Lasso
- This is currently the most popular version of regularized regression

IMPLEMENTATION

```
set.seed(123)
```

```
cv.lasso <- cv.glmnet(  
  x = train.x,  
  y = as.factor(train.y),  
  alpha = 1,  
  family = "binomial",  
  nfolds = 10,  
  type.measure = "class"  
)
```

`set.seed` provides reproducibility

IMPLEMENTATION

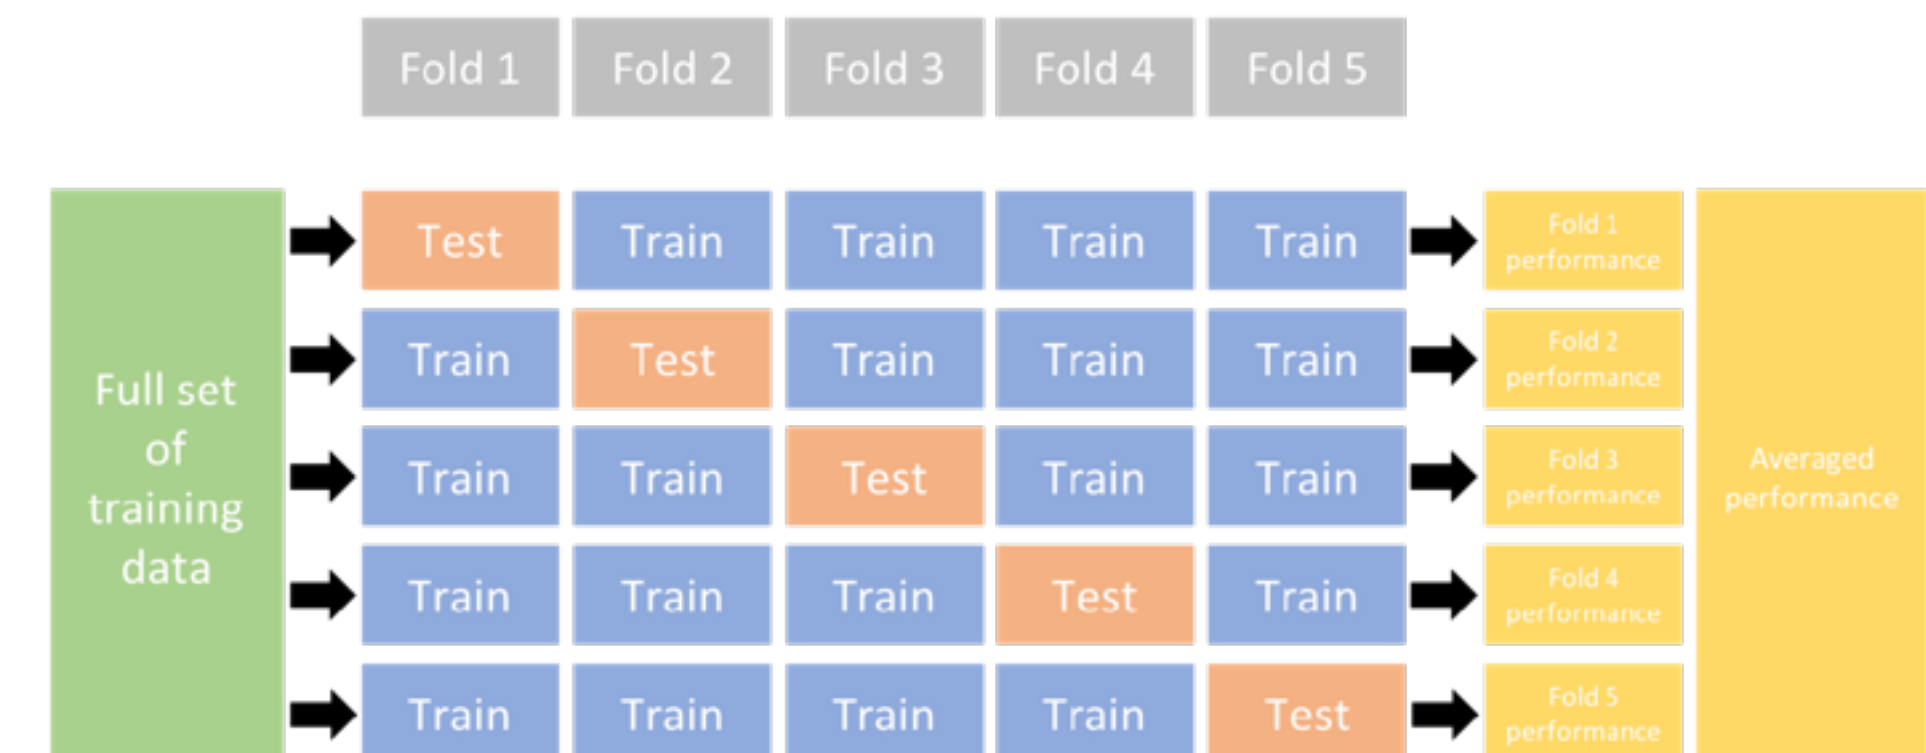
```
set.seed(123)

cv.lasso <- cv.glmnet(
  x = train.x,
  y = as.factor(train.y),
  alpha = 1,
  family = "binomial",
  nfolds = 10,
  type.measure = "class"
)
```

`cv.glmnet` provides built-in cross validation.

There is an equivalent `glmnet` function w/o built-in cv.

We use `nfolds` to tell the model how many cv folds to apply.



IMPLEMENTATION

```
set.seed(123)

cv.lasso <- cv.glmnet(
  x = train.x,
  y = as.factor(train.y),
  alpha = 1,
  family = "binomial",
  nfolds = 10,
  type.measure = "class"
)
```

`x` & `y` feed in our training and response data.

When performing classification we want our response to be encoded **`as.factor`**

IMPLEMENTATION

```
set.seed(123)

cv.lasso <- cv.glmnet(
  x = train.x,
  y = as.factor(train.y),
  alpha = 1,
  family = "binomial",
  nfolds = 10,
  type.measure = "class"
)
```

alpha is the parameter that tells it which of the regularization methods to use:

- **alpha = 1**: Lasso
- **alpha = 0**: Ridge
- **0 < alpha < 1**: Elastic net

IMPLEMENTATION

```
set.seed(123)

cv.lasso <- cv.glmnet(
  x = train.x,
  y = as.factor(train.y),
  alpha = 1,
  family = "binomial",
  nfolds = 10,
  type.measure = "class"
)
```

family describes the distribution of the y response variable.

Many to choose from but most common:

- **gaussian**: continuous
- **binomial**: binomial for two level classification


IMPLEMENTATION

```
set.seed(123)

cv.lasso <- cv.glmnet(
  x = train.x,
  y = as.factor(train.y),
  alpha = 1,
  family = "binomial",
  nfolds = 10,
  type.measure = "class"
)
```

`type.measure` states the performance measure used during the cv process:

- `class`: minimize misclassification rate
- `AUC`: maximize area under the curve
- `MSE`: minimize mean squared error
- `class`: minimize absolute square error




ASSESSMENT

```
set.seed(123)

cv.lasso <- cv.glmnet(
  x = train.x,
  y = as.factor(train.y),
  alpha = 1,
  family = "binomial",
  nfolds = 10,
  type.measure = "class"
)

plot(cv.lasso)
```

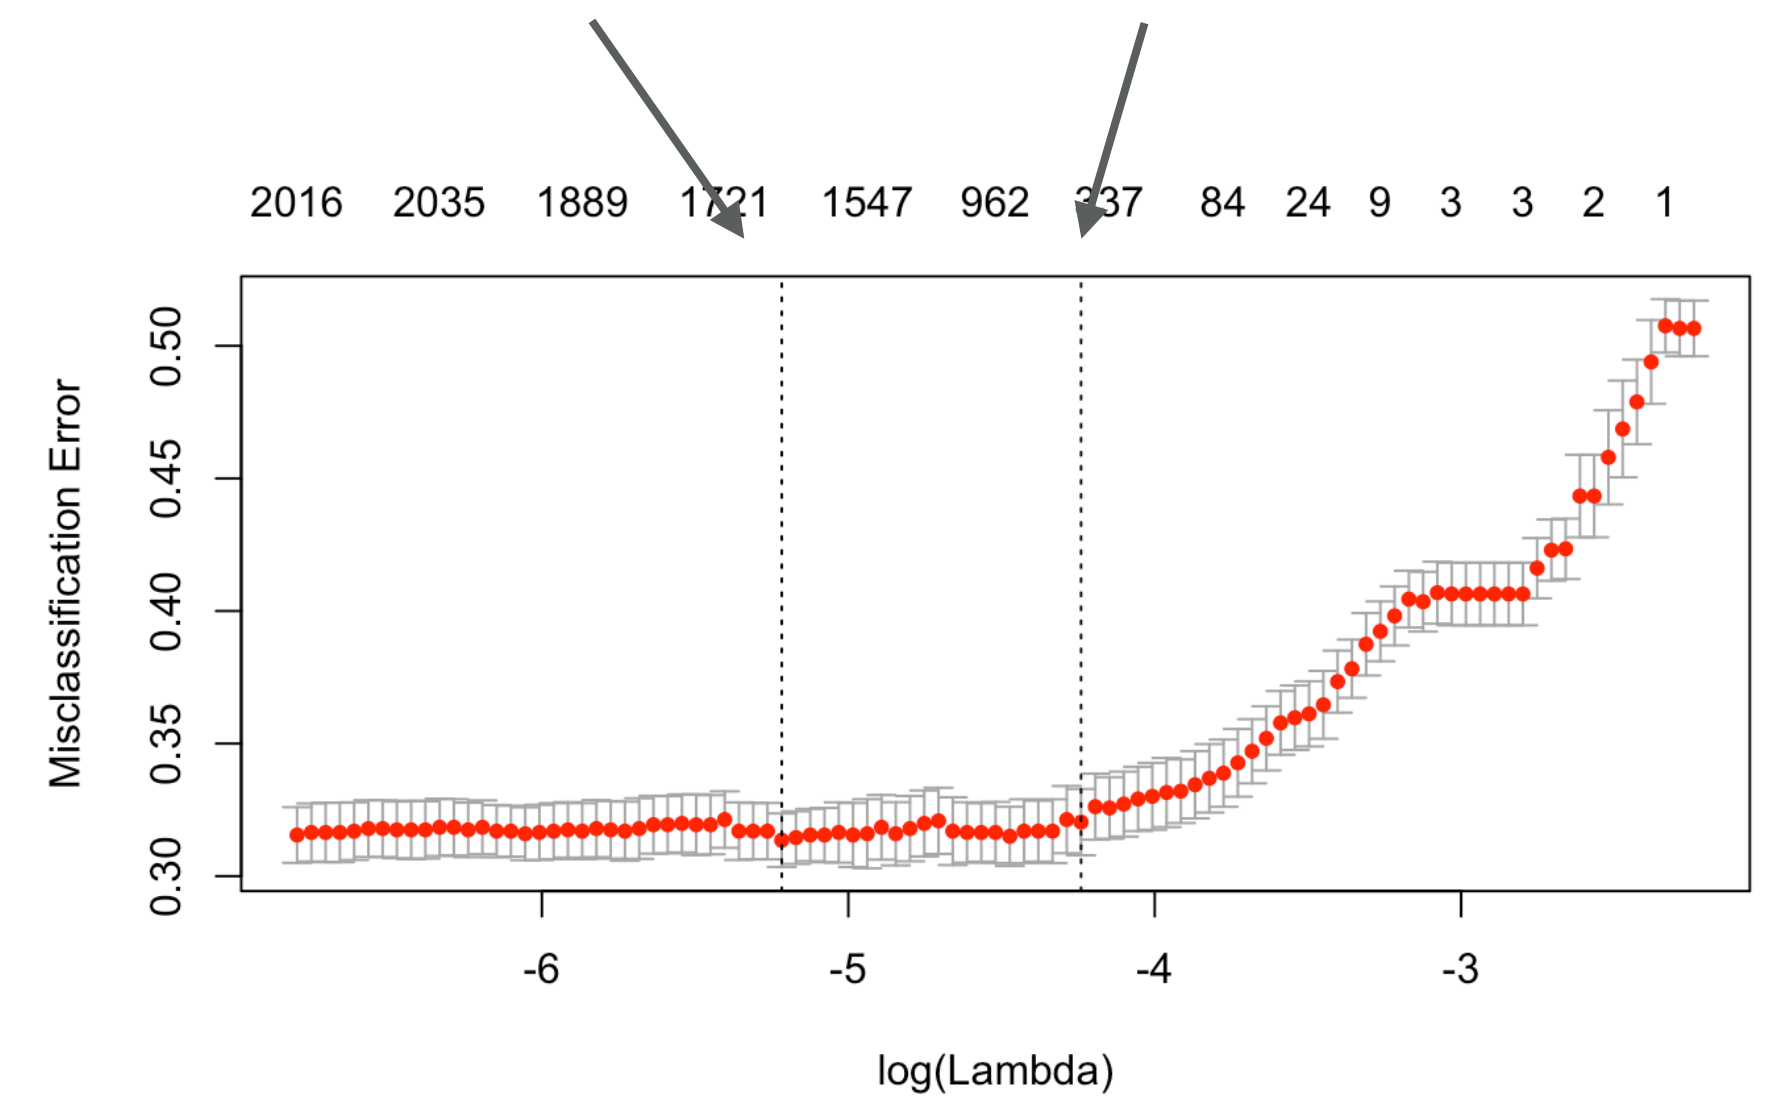


ASSESSMENT

```
# predict lasso
pred.lasso <- predict(
  cv.lasso,
  train.x,
  type = "class",
  s = cv.lasso$lambda.1se
)
```

To get predictions we feed predict:

- `cv.lass`: model
- `train.x`: new feature set
- `type`: prediction type (class or probability)
- `s`: `lambda.min` or `lambda.1se`



ASSESSMENT


```
# predict lasso
pred.lasso <- predict(
  cv.lasso,
  train.x,
  type = "class",
  s = cv.lasso$lambda.1se
)

auc.lasso <- roc(train.y, as.numeric(pred.lasso))
auc.lasso
Area under the curve: 0.8359

plot(auc.lasso)
```

We can assess the area under the curve with **pROC::roc**

- Provides us with how well the model balances:
 - True positives
 - True negatives



ASSESSMENT

```
# confusion matrix
confusion <- table(pred.lasso, train.y)
confusion
      train.y
pred.lasso  0    1
      0 731   33
      1 307  986

# overall accuracy
sum(diag(confusion)) / sum(confusion)
[1] 0.8347107
```

We can also the confusion matrix

YOUR TURN I!

Apply a Ridge regression model (hint: $\alpha = ?$) to the same training data and plot the misclassification rate. Which range of lambdas performs the best?

YOUR TURN 2!

*Now get the predictions for the `train.x` data using `s = cv.ridge$lambda.1se`.
What is the area under the curve? Can you create the confusion matrix?*

ELASTIC NET

Let's assess how an elastic net compares

We can apply an elastic net model by choosing $0 < \alpha < 1$

But what value do we pick?

ELASTIC NET

```
# Let's assess how an elastic net compares
tuning <- expand.grid(
  alpha = seq(0, 1, by = .01),
  accuracy = NA
)
```

```
head(tuning)
  alpha accuracy
1  0.00      NA
2  0.01      NA
3  0.02      NA
4  0.03      NA
5  0.04      NA
6  0.05      NA
```

We can apply an elastic net model by choosing $0 < \alpha < 1$

I. Let's create a tuning grid

ELASTIC NET

```
# Let's assess how an elastic net compares
for(i in seq_along(tuning$alpha)) {

  set.seed(123)

  cv <- cv.glmnet(
    x = train.x,
    y = as.factor(train.y),
    alpha = tuning$alpha[i],
    family = "binomial",
    nfolds = 10,
    type.measure = "class"
  )


  pred <- predict(cv, train.x, type = "class", s = cv$lambda.1se)
  confusion <- table(pred, train.y)
  tuning$accuracy[i] <- sum(diag(confusion)) / sum(confusion)
}
```

We can apply an elastic net model by choosing $0 < \alpha < 1$

1. Let's create a tuning grid
2. Now we'll loop through each alpha value, apply a glmnet model, compute the accuracy and add that value into our tuning grid.

ELASTIC NET

```
# Let's assess how an elastic net compares  
ggplot(tuning, aes(alpha, accuracy)) +  
  geom_line() +  
  ylim(c(0, 1))
```



We can apply an elastic net model by choosing $0 < \alpha < 1$

1. Let's create a tuning grid
2. Now we'll loop through each alpha value, apply a glmnet model, compute the accuracy and add that value into our tuning grid.
3. Now we can plot our results across all alpha values.

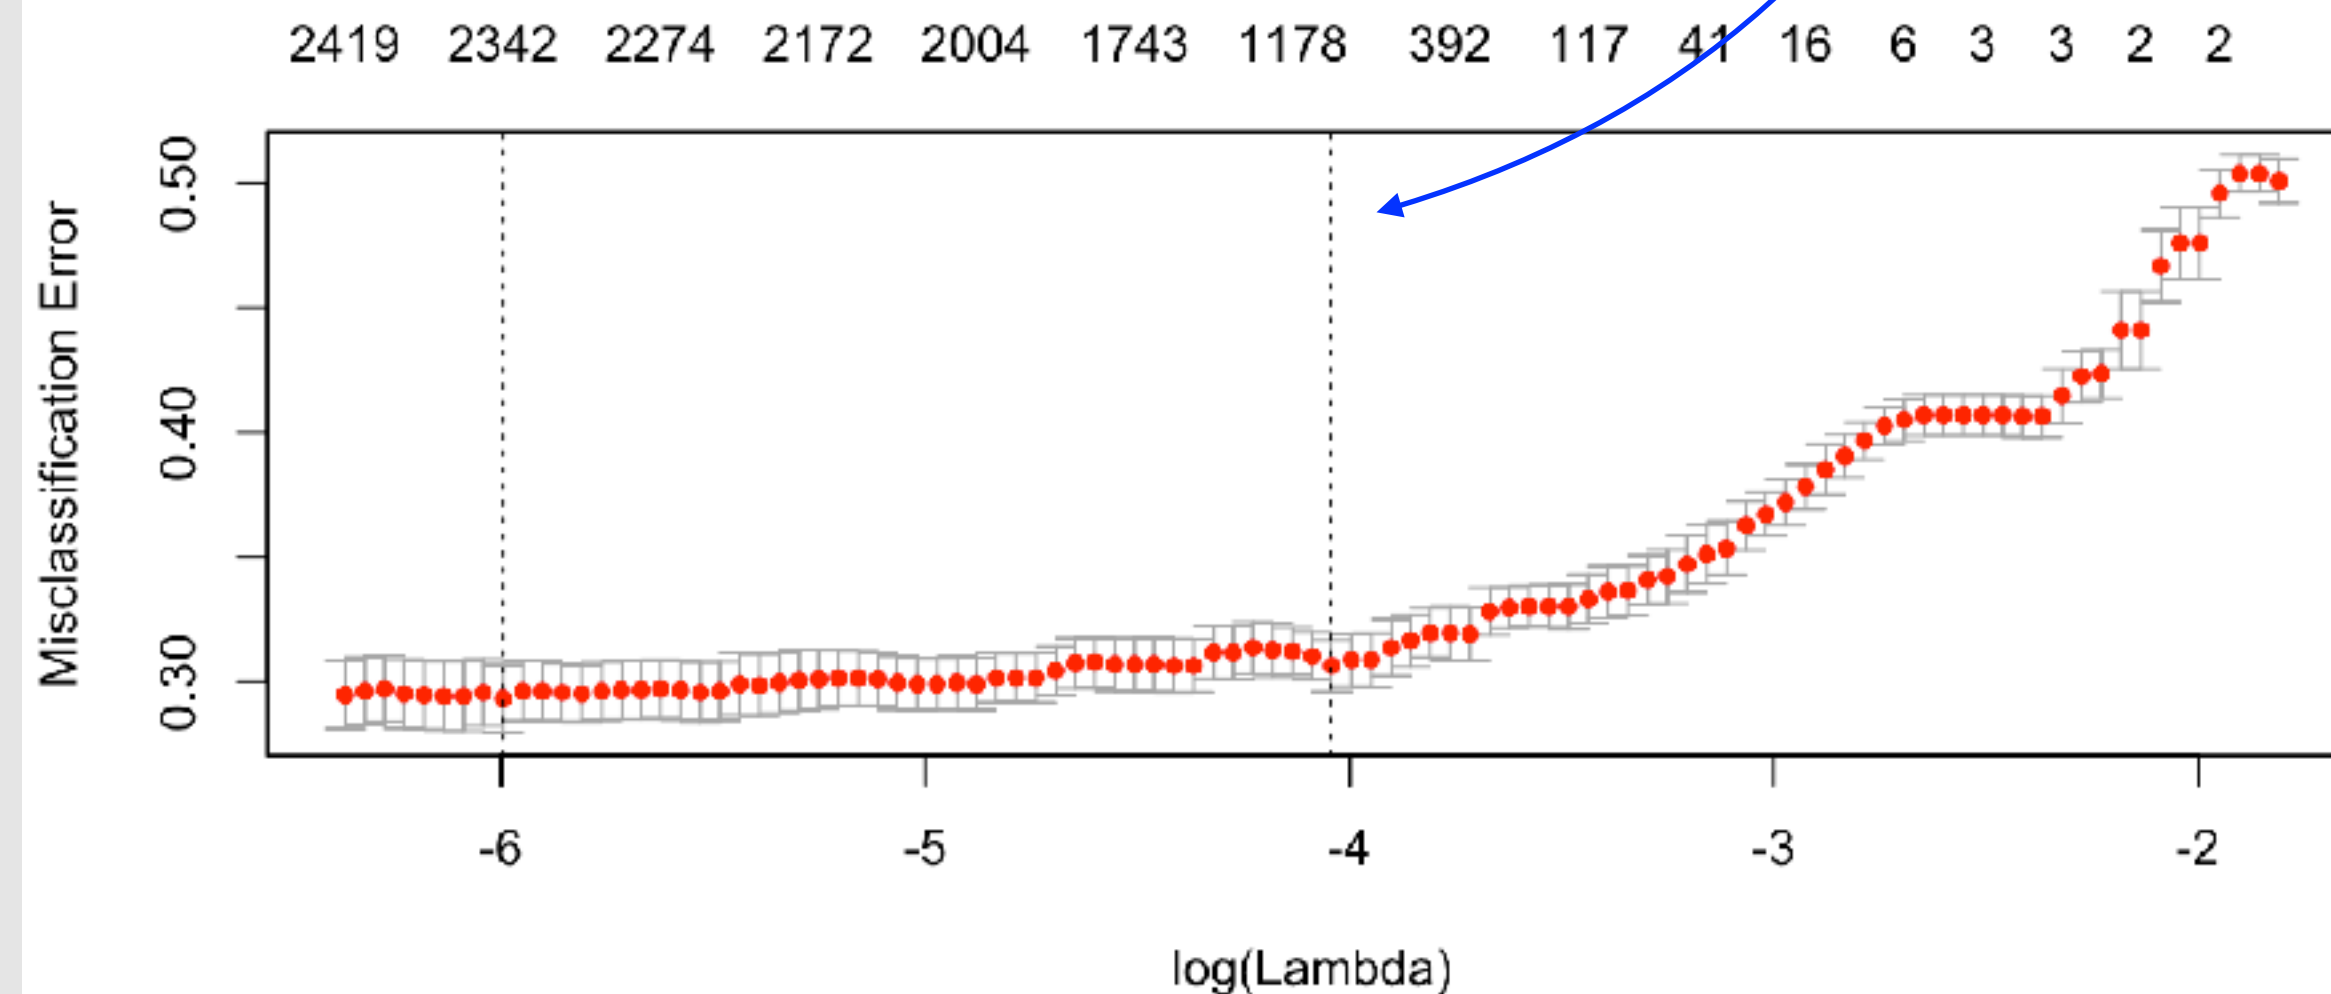
ELASTIC NET

```
# let's apply an elastic net with alpha = .65
```

```
cv <- cv.glmnet(  
  x = train.x,  
  y = as.factor(train.y),  
  alpha = .65,  
  family = "binomial",  
  nfolds = 10,  
  type.measure = "class"  
)
```

Let's apply an elastic net with
alpha = .65

We can reduce our feature set down to about 1000 if we wanted/needed to.



ELASTIC NET

```
# Now let's predict...but with our test set
```

```
pred.test <- predict(
```

```
  cv,
```

```
  test.x,
```

```
  type = "class",
```

```
  s = cv$lambda.min
```

```
)
```

```
auc <- roc(test.y, as.numeric(pred.test))
```


```
auc
```

```
Data: as.numeric(pred.test) in 450 controls (test.y 0) <  
440 cases (test.y 1).
```

```
Area under the curve: 0.7178
```


```
plot(auc)
```

Now let's assess how well our model generalizes to an unseen data set (**test.x**)



ELASTIC NET

Grey solid = Ridge, Grey dotted = Lasso, Red = Elastic



UNDERSTANDING THE RESULTS

So we found a decent model, now
how do we understand the
results?

UNDERSTANDING THE RESULTS

```
# finding the most impactful words
glmnet.coef <- as.matrix(coef(cv, s = "lambda.min"))
glmnet.coef <- tibble(
  words = row.names(glmnet.coef),
  coeff = glmnet.coef[, 1]
) %>%
  arrange(desc(coeff)) %>%
  mutate(words = fct_inorder(words))
```

glmnet.coef

A tibble: 7,150 x 2

	words	coeff
	<fct>	<dbl>
1	overthrow	6.47
2	pummel	5.55
3	u.k	4.54
4	doug	4.54

We can extract each feature (word) and the coefficient

UNDERSTANDING THE RESULTS


```
# finding the most impactful words
```

```
summary(glmnet.coef$coeff)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
-5.581671	0.000000	0.000000	0.006956	0.000000	6.472727

```
ggplot(glmnet.coef, aes(coeff)) +  
  geom_density()
```


We see that many of the features have coefficients that have been shrunk to zero



UNDERSTANDING THE RESULTS

```
# top 10 positive and negative words
rbind(
  top_n(glmnet.coef, 10),
  top_n(glmnet.coef, -10)
) %>%
  ggplot(aes(x = coeff, y = words)) +
  geom_segment(aes(yend = words, xend = 0)) +
  geom_point(aes(color = coeff > 0), size = 2)
```

We can look at the top 20 influential words



UNDERSTANDING THE RESULTS

```
# convert to probabilities
glmnet_coef <- glmnet.coef %>%
  mutate(
    probability = arm::invlogit(coeff),
    word_id     = row_number(words)
  )
```

glmnet_coef

A tibble: 7,150 x 4


	words	coeff	probability	word_id
	<fct>	<dbl>	<dbl>	<int>
1	overthrow	6.47	0.998	1
2	pummel	5.55	0.996	2
3	u.k	4.54	0.989	3
4	doug	4.54	0.989	4
5	brazilian	4.29	0.986	5
6	patience	4.15	0.984	6

Let's compute the probabilities

UNDERSTANDING THE RESULTS

```
# plot word probability distributions  
ggplot(glmnet_coef, aes(probability, word_id)) +  
  geom_point(size = .5)
```




Let's visualize the probabilities



SO LITTLE TIME!



LEARN MORE



WHAT TO REMEMBER



FUNCTIONS TO REMEMBER

Operator/Function	Description
<code>Matrix::Matrix</code>	Create a sparse matrix
<code>glmnet::cv.glmnet</code>	Compute a cross validated regularized regression model (Ridge: <code>alpha = 0</code> ; Lasso: <code>alpha = 1</code> ; Elastic net: <code>0 < alpha < 1</code>)
<code>pROC::roc</code>	Create a ROC curve and compute area under the curve
<code>predict</code>	Predict response values for a new feature set